
Customer Lifetime Value Prediction

Using Bayesian Methods

A Probabilistic Approach to Non-Contractual
Customer-Base Analysis in E-Commerce

Master's Thesis

submitted in partial fulfilment of the requirements
for the degree of Master of Science

Devansh

June 15, 2026

Abstract

Customer Lifetime Value (CLV) is the expected net present value of all future cash flows from a customer relationship and is a foundational input to marketing-resource allocation. In non-contractual settings such as e-commerce, customers never announce their departure, so the central modelling difficulty is probabilistic reasoning about an unobservable attrition state. This thesis develops and evaluates a Bayesian probabilistic approach to CLV prediction built on the Beta-Geometric/Negative-Binomial-Distribution (BG/NBD) model for purchase frequency and the Gamma-Gamma model for monetary value, estimated by Hamiltonian Monte Carlo with the No-U-Turn Sampler. Three contributions are examined empirically on the UCI Online Retail II dataset: (i) whether the Bayesian model matches or exceeds RFM and XGBoost baselines on out-of-sample accuracy while additionally providing calibrated uncertainty (RQ1/H1); (ii) whether a hierarchical extension with partial pooling across country segments improves prediction for data-sparse markets (RQ2/H2); and (iii) whether decision-theoretic targeting based on the posterior probability $P(\text{CLV} > c)$ realises greater value than targeting by point-estimate CLV (RQ3/H3). The theoretical chapters reproduced here establish the conceptual landscape, the generative models, the inferential machinery, and the three testable hypotheses that the empirical analysis evaluates.

Contents

Abstract

1	Introduction	1
1.1	Background and Motivation	1
1.2	Understanding Customer Lifetime Value: Concept and Historical Development	1
1.3	Customer Lifetime Value in E-Commerce and Modern Business	3
1.4	Research Questions and Hypotheses	4
1.5	Structure of the Thesis	4
2	Theoretical Foundations of Customer Lifetime Value	6
2.1	Customer Lifetime Value: Concept	6
2.1.1	Contractual vs. Non-Contractual Settings	6
2.2	Pre-Data and Non-Data Approaches to CLV Estimation	7
2.3	Traditional Approaches to CLV Estimation	8
2.3.1	RFM Scoring and Heuristic Segmentation	8
2.3.2	Deterministic CLV Formulas	9
2.4	Machine-Learning Approaches to CLV	11
2.5	Summary of Theoretical Findings	12
3	Bayesian Probabilistic CLV Models	13
3.1	Bayesian Inference Fundamentals	13
3.1.1	Markov Chain Monte Carlo and the NUTS Sampler	14
3.2	The BG/NBD Model for Purchase Frequency	14
3.2.1	Generative Story	14
3.2.2	Key Model Outputs	16
3.2.3	Model Assumptions and Limitations	17
3.3	The Gamma-Gamma Model for Monetary Value	17
3.3.1	Model Specification	17

3.3.2	Key Assumption: Independence of Frequency and Monetary Value .	18
3.3.3	CLV Computation	18
3.4	Hierarchical Extensions: Partial Pooling Across Segments	19
3.4.1	Partial Pooling vs. Alternatives	19
3.4.2	Implementation Considerations	21
3.5	Decision Theory Under Uncertainty	21
3.5.1	The Targeting Problem	21
3.5.2	Expected Improvement Over Point Estimates	22
3.6	Summary of Theoretical Findings and Hypotheses	22
4	Methodology	24
4.1	Research Design	24
4.2	Dataset	24
4.3	Data Preparation	25
4.3.1	Cleaning	25
4.3.2	Temporal Split	25
4.3.3	Customer-Level Aggregation	25
4.3.4	Country Segmentation	27
4.4	Exploratory Findings	27
4.5	Model Specification	28
4.6	Bayesian Inference	29
4.7	Baseline Models	29
4.8	Prediction	29
4.9	Evaluation Protocol	30
4.10	Implementation and Reproducibility	31
5	Results	32
5.1	MCMC Convergence and Diagnostics	32
5.2	H1: Predictive Accuracy and Calibrated Uncertainty	34
5.2.1	Point and Ranking Accuracy	34

5.2.2	Uncertainty Calibration	34
5.2.3	Activity Prediction	35
5.3	H2: Hierarchical Partial Pooling Across Segments	36
5.4	H3: Decision-Theoretic Targeting	37
5.5	Summary of Hypothesis Tests	38
6	Discussion and Conclusion	39
6.1	Summary of Findings	39
6.2	Theoretical Contributions	40
6.3	Managerial Implications	40
6.4	Limitations	40
6.5	Future Work	41
6.6	Concluding Remarks	41
	References	42

List of Figures

3.1	Graphical model of the BG/NBD + Gamma-Gamma framework	15
3.2	Buy-till-you-die timeline	16
3.3	Empirical recency versus observation period (T) for the calibration sample. Every customer necessarily lies on or below the diagonal recency = T ; the spread of points away from the diagonal encodes the latent-activity signal the BG/NBD model exploits.	16
3.4	Correlation structure among the RFM summary statistics in the calibration data. The frequency–monetary association is the quantity that must be near zero for the Gamma-Gamma independence assumption to hold.	18
3.5	Hierarchical partial pooling	19
3.6	Customer counts by country segment in the calibration data. The pronounced imbalance, a dominant UK segment alongside several small markets, is precisely the regime in which partial pooling is expected to help.	20
3.7	Decision theory under uncertainty	22
4.1	Data-cleaning waterfall: the cumulative effect of each filtering step on the row count, from 1,067,371 raw line items to 802,679 retained transactions. .	25
4.2	Monthly transaction revenue across the full observation window. The calibration/holdout boundary is fixed at 1 March 2011.	26
4.3	Distribution of repeat-purchase frequency in the calibration sample. The mass at low frequencies and the long right tail are characteristic of non-contractual retail.	27
4.4	Distribution of average transaction value among repeat purchasers. The right skew motivates the Gamma-Gamma specification.	28
5.1	Split- $\hat{R}$ convergence diagnostics for the pooled BG/NBD model. The dashed line marks the 1.01 threshold.	32
5.2	Posterior trace plots for the pooled BG/NBD population parameters (r, α, a, b) . Well-mixed, stationary chains indicate reliable sampling.	33
5.3	Pairwise posterior marginals for the BG/NBD parameters, with divergences highlighted. Strong correlations would motivate reparameterisation.	33

5.4	Split- $\hat{R}$ for the hierarchical BG/NBD model under the non-centred parameterisation.	33
5.5	Decile calibration: mean predicted versus mean actual holdout transactions for each model. The diagonal denotes perfect calibration.	34
5.6	Posterior predictive distributions of holdout transactions for a sample of customers (violins), with realised values overlaid. The spread of each violin is the model's per-customer uncertainty.	35
5.7	Distribution of $P(\text{alive})$ overall and by frequency group. Higher-frequency customers concentrate at higher $P(\text{alive})$	36
5.8	Per-country transaction MAE by model. Differences concentrated in the small segments (Germany, France, Other) test the partial-pooling hypothesis.	36
5.9	Per-segment posterior of the BG/NBD parameter r (90% HDI), with the global mean marked. Small segments are shrunk toward the global estimate.	37
5.10	Decision-theoretic targeting for the Bayesian model at the headline cost of £20. Left: cumulative net value by strategy and depth. Right: the net-value advantage of the posterior rule over the point estimate.	37
5.11	Cumulative-gain and lift curves for all models. Steeper early rise indicates better identification of high-value customers.	38

List of Tables

2.1	The evolution of CLV estimation, from deterministic heuristics to generative Bayesian models. Only the probabilistic-generative tradition produces calibrated, distributional value estimates from raw transaction data.	10
3.1	Estimation strategies for multi-segment customer data.	20
3.2	Testable hypotheses evaluated in the empirical analysis.	23
4.1	Mapping of evaluation metrics to research hypotheses.	30
5.1	Model comparison across transaction and CLV prediction metrics. [pending pipeline run]	34
5.2	Posterior predictive credible-interval coverage and CRPS. [pending pipeline run]	35
5.3	$P(\text{alive})$ evaluation: binary activity-prediction metrics. [pending pipeline run]	35
5.4	Per-country transaction-prediction MAE by model (RQ2/H2). [pending pipeline run]	36
5.5	Decision-theoretic targeting (BG/NBD): posterior-probability versus point-estimate versus oracle net value, swept over intervention cost. [pending pipeline run]	37
5.6	Summary of hypothesis-test outcomes.	38

Chapter 1

Introduction

1.1 Background and Motivation

In the modern digital economy, the ability to anticipate future customer behaviour is among the most commercially valuable capabilities a business can possess. Whether allocating limited marketing budgets, designing retention programmes, or pricing acquisition channels, managers routinely face decisions whose quality depends on how well they can estimate the long-term value of individual customers. Customer Lifetime Value (CLV), the expected net present value of all future cash flows from a customer relationship, provides a theoretically grounded and practically actionable foundation for these decisions (Gupta et al., 2006). As Blattberg and Deighton (1996) argued, customers are financial assets whose value can and should be measured with the same rigour applied to capital investments; CLV formalises this intuition into a computable quantity.

Despite its conceptual clarity, accurately estimating CLV in practice is considerably more complex than the underlying formula implies. In non-contractual contexts, which are prevalent in e-commerce, where customers transact sporadically without formal subscriptions, organisations cannot directly ascertain when a customer has permanently churned. At any given time, a customer who has not transacted for several months may simply be in an extended inter-purchase interval or may have silently defected. This latent attrition issue implies that naive extrapolation of historical purchase frequencies will systematically overestimate the value of customers who have, in reality, discontinued their relationship. Consequently, the central challenge extends beyond predictive accuracy to encompass probabilistic reasoning under uncertainty about an inherently unobservable state (Fader & Hardie, 2009).

1.2 Understanding Customer Lifetime Value: Concept and Historical Development

At its essence, CLV is a financial construct: it denotes the anticipated net present value of all future cash flows a firm expects from a customer relationship, adjusted to reflect their present value. Rather than focusing solely on isolated transactions, CLV encourages a forward-looking viewpoint. It prompts firms to consider not only what a customer has previously spent, but also to estimate their future spending potential, the likely duration of their engagement, and the associated degree of uncertainty. This reframing influences numerous

managerial decisions, including the allocation of marketing budgets, the structure of retention initiatives, and the broader assessment of customer-base health. Blattberg and Deighton (1996) advanced the notion that customers should be treated as financial assets, deserving of valuation methods as robust as those used for capital investments. The sum of these individual values, customer equity, forms a critical element of total firm value and merits careful management. Empirical evidence from Gupta et al. (2006) underscores this perspective, revealing that even modest gains in customer retention can yield disproportionately large increases in overall firm value. Such sensitivity highlights the commercial significance of refining CLV-model accuracy, even incrementally. In practice, accurate estimation of individual customer value allows firms to prioritise acquisition efforts, target retention spending more effectively, and reduce resource allocation to relationships unlikely to justify further investment.

The development of CLV as a concept mirrors a broader transformation in marketing, marking a shift from a transaction-based orientation, where each sale was viewed as an independent occurrence and effectiveness was judged by aggregate indicators such as market share and revenue, toward a relationship-centric approach. In this newer paradigm, the enduring profitability of individual customer relationships assumes central importance. Research has highlighted the shortcomings of the transaction model, particularly as evidence mounted that customer value is not uniformly distributed but instead exhibits marked heterogeneity: a relatively small subset of customers often accounts for a significant share of overall revenue, a pattern Storbacka (1997) articulated through the customer-portfolio lens. The urgency of managing this heterogeneity was reinforced by Reichheld and Sasser (1990), whose findings demonstrated that even marginal reductions in customer defection rates could yield substantial profitability improvements. Their work encouraged firms to track and enhance individual customer value over time, moving beyond a one-size-fits-all retention approach. The RFM (Recency, Frequency, Monetary value) framework subsequently gained traction among practitioners for its clarity and ease of use, yet its retrospective focus limited its capacity to predict future behaviour or provide probabilistic value estimates (Hughes, 1994). A more robust theoretical foundation was set by Dwyer (1989), who positioned CLV as a net-present-value calculation, and further advanced by the emergence of probabilistic “buy-till-you-die” models (Schmittlein et al., 1987). These models treated purchasing and attrition as stochastic processes, enabling firms to infer the probability of a customer remaining active from observed transactions. Fader et al. (2005) later developed the BG/NBD model, simplifying earlier probabilistic frameworks while retaining their core logic. More recently, advances in machine learning have introduced discriminative models trained on rich behavioural data to forecast future spending, though these methods have drawn criticism for their sometimes poor uncertainty calibration and limited generalisability (Chamberlain et al., 2017). The evolution of CLV modelling therefore spans deterministic methods, probabilistic generative models, and machine-learning approaches, each addressing unique challenges

in estimation. This thesis positions itself within the probabilistic tradition, contending that Bayesian inference applied to generative CLV models offers a particularly rigorous framework for estimating individual customer value under uncertainty, and seeks to validate this assertion empirically.

1.3 Customer Lifetime Value in E-Commerce and Modern Business

The significance of CLV has reached new heights in the current e-commerce landscape, where the economics of digital retail make precise measurement at the customer level not merely advantageous but essential. In contrast to traditional brick-and-mortar stores, where acquisition costs are spread across physical infrastructure and broad brand presence, online retailers face explicit and quantifiable per-customer acquisition expenses, ranging from paid search and affiliate fees to social-media campaigns and influencer collaborations. These costs must be justified by the anticipated long-term returns from each customer. In many e-commerce sectors, initial purchases rarely cover acquisition expenses, making accurate projections of future customer value vital to sustainable business models. Overestimating CLV leads to excessive bidding for new customers and value destruction, whereas underestimating it results in missed opportunities and loss of market share to more aggressive competitors. Gupta et al. (2006) highlight this tension as a core strategic consideration within CLV modelling, a concern now amplified by rising digital-marketing costs and heightened competition across nearly all product categories.

CLV serves as a critical analytical tool for a spectrum of strategic decisions, especially within e-commerce. For effective retention marketing, encompassing discount strategies, loyalty programmes, personalised recommendations, and efforts to re-engage inactive customers, a robust estimate of future customer value is indispensable. Without such forward-looking metrics, firms struggle to assess whether the cost of retention initiatives is justified by potential gains, or to distinguish between customers meriting resource-intensive interventions and those for whom automated, low-cost outreach suffices. E-commerce businesses benefit from a wealth of behavioural data, including clickstreams, browsing patterns, cart-abandonment rates, and cross-category purchasing behaviour. However, the sheer volume and complexity of this data introduce new analytical hurdles: relying solely on past behaviour without a sound model can skew predictions, particularly for high-value or low-frequency customers who are pivotal to targeting strategies. Reinartz and Kumar (2000) provided early empirical evidence that the link between customer tenure and profitability is far more nuanced than widely assumed: many long-term customers are not profitable, while some short-term customers contribute outsized value, underscoring the need for individual-level CLV models rather than broad retention benchmarks.

The emergence of subscription-based and platform-driven business models has made CLV

an even more pivotal metric for strategic decision-making. McCarthy and Fader (2018) illustrated that valuing a company by aggregating individual customer lifetime values across its user base often yields more precise and timely firm valuations than conventional accounting methods, with important ramifications for investor communications, acquisitions, and long-term planning. In e-commerce, where firms may operate at a loss while aggressively acquiring new customers, CLV-based analyses allow stakeholders to distinguish unsustainable, value-eroding growth from rational investment in customer assets expected to generate future returns. Such analytical distinctions, often overlooked by traditional financial reporting, have become central to how informed investors and executives assess business viability. For the purposes of this thesis, the UCI Online Retail II data under review exemplifies a non-contractual, episodic purchasing setting, precisely the environment in which robust CLV estimation is both challenging and impactful, and where the Bayesian probabilistic methods employed here offer clear advantages over more basic approaches.

1.4 Research Questions and Hypotheses

This thesis is organised around three research questions (RQ) and their associated hypotheses (H), each evaluated empirically in the results chapter and motivated theoretically in Chapter 3:

RQ1 / H1: Accuracy and calibrated uncertainty.

Does the Bayesian BG/NBD + Gamma-Gamma model achieve competitive or superior out-of-sample accuracy relative to RFM-heuristic and XGBoost baselines, while additionally delivering well-calibrated uncertainty estimates?

RQ2 / H2: Hierarchical partial pooling.

Does a hierarchical extension that partially pools information across country segments reduce prediction error for small-market customers relative to complete-pooling and no-pooling alternatives?

RQ3 / H3: Decision-theoretic targeting.

Does targeting customers by the posterior probability $P(\text{CLV} > c)$ produce higher cumulative realised value than targeting by point-estimate CLV rankings?

1.5 Structure of the Thesis

Chapter 2 reviews the theoretical foundations of CLV estimation, tracing the progression from pre-data heuristics through deterministic formulas and RFM scoring to modern machine-learning approaches, and isolating the gaps that motivate a probabilistic treatment. Chapter 3 develops the Bayesian probabilistic framework that constitutes the methodological core of

the thesis: Bayesian inference and MCMC, the BG/NBD and Gamma-Gamma generative models, their hierarchical extension via partial pooling, and the decision-theoretic targeting framework. The chapter closes by formalising the three hypotheses above.

Chapter 2

Theoretical Foundations of Customer Lifetime Value

2.1 Customer Lifetime Value: Concept

CLV represents the total net present value of all future cash flows attributed to a customer over the entirety of their relationship with a firm. As a forward-looking metric, CLV shifts managerial attention from short-term transactional revenue toward long-term customer equity, enabling firms to make economically rational decisions about acquisition spending, retention investment, and resource allocation across their customer base (Gupta et al., 2006).

The concept of CLV emerged from the broader customer-equity literature pioneered by Blattberg and Deighton (1996), who argued that customers should be treated as financial assets whose value can be measured, managed, and optimised. Rust et al. (2000) formalised this perspective into the customer-equity framework, decomposing total firm value into the sum of individual customer lifetime values. This conceptual shift has profound implications for marketing practice: rather than evaluating campaigns solely on immediate return, marketers can assess investments based on their impact on the long-term value of the customer base. Formally, CLV can be expressed as

$$\text{CLV} = \sum_{t=1}^T \frac{m_t r_t}{(1+d)^t}, \quad (2.1)$$

where m_t is the expected margin in period t , r_t is the retention probability, d is the discount rate, and T is the time horizon. This deterministic formulation, while intuitive, requires point estimates of future margins and retention rates, introducing substantial uncertainty that the formula itself does not capture. This limitation motivates the probabilistic approaches that form the core of this thesis.

2.1.1 Contractual vs. Non-Contractual Settings

A critical distinction in the CLV literature is between contractual and non-contractual business settings (Fader & Hardie, 2009). In contractual settings (subscription services, insurance policies, or mobile-phone contracts), the firm observes when a customer becomes inactive, because the customer explicitly cancels or fails to renew. The modelling task reduces to predicting *when* customers will terminate, and standard survival-analysis methods apply directly.

Non-contractual settings, typical of most retail and e-commerce environments, present a fundamentally harder problem. Customers do not announce their departure; they simply stop purchasing. At any point, the firm cannot distinguish with certainty between a customer who has permanently defected and one who is merely experiencing a long inter-purchase interval. This latent-attribution problem requires probabilistic models that jointly estimate the customer's transaction rate and their unobserved dropout process, making it a natural domain for Bayesian inference. This thesis focuses exclusively on the non-contractual setting, as it represents the more challenging and practically common scenario in e-commerce. The UCI Online Retail II dataset used in the empirical part reflects this setting: customers purchase episodically with no formal contract, and the firm must infer activity status from observed purchase patterns.

2.2 Pre-Data and Non-Data Approaches to CLV Estimation

Prior to the advent of sophisticated transaction-level data and the computational power required to analyse it, organisations relied on a combination of judgment-based and formulaic strategies to estimate CLV. Although these early approaches lacked the precision of contemporary models, they laid the groundwork for the conceptual frameworks that underpin current CLV methodologies. The most prevalent legacy technique was the deterministic CLV formula, which calculated expected customer value by multiplying average order value, purchase frequency, and an estimated customer lifespan, then discounting to present value using a managerially determined rate. This method, detailed in the direct-marketing literature by Hughes (1994) and Stone (1995), was valued for its clarity: managers could trace the logic from inputs to outcome, which made it effective for justifying budgets or informing strategic discussions. Its core drawback lay in its reliance on assumed or historically averaged inputs, offering no means to capture genuine differences among individual customers, precisely the nuance needed for effective targeting. Applying the same expected lifespan to both a frequent purchaser and a one-time buyer, with only average order value as a differentiator, conflated distinct customer types and led to aggregate estimates that might be reasonable in total but largely uninformative at the individual level.

A related method, cohort-based value estimation, grouped customers acquired in the same timeframe and tracked their combined revenue over subsequent periods, producing a survival curve that firms could extrapolate to predict the remaining value of active members. Dwyer (1989) identified this as one of the earliest formal CLV treatments. While it leveraged actual retention data, improving on the single-formula method, the approach still operated at an aggregate level and required untestable assumptions about the projected survival curve beyond observed data.

An incremental move toward individualisation came with RFM scoring, which derived com-

posite scores from customers' transaction histories, allowing companies to identify high-value customers for targeted marketing without building formal statistical models. Popularised by Hughes (1994) and Cullinan (1977) in direct mail, RFM remains widespread in practice. Yet its limitations are clear: it reflects only past behaviour, lacks any measure of confidence or predictive reliability for individual rankings, and cannot distinguish genuinely promising customers from those who simply made an unusually large recent purchase.

Expert judgment and managerial intuition represented another tier of pre-data CLV estimation, especially in business-to-business contexts where transaction volumes were low and individual insights could be obtained directly by account teams. Sales teams' estimates of renewal probabilities and expected contract values served as informal proxies for CLV, which management aggregated into revenue forecasts. While subjective, these estimates often incorporated relationship-specific intelligence inaccessible to algorithmic models. However, Lovallo and Kahneman (2003) and others document systematic biases (overconfidence in renewals, recency effects, and the inflation of estimates to protect internal resources) that ultimately motivated the transition to data-driven CLV models as more granular data became available.

A unifying feature of all these pre-data approaches is their reduction of uncertainty to a fixed parameter rather than an estimable quantity. Each method delivered a single point estimate per customer or cohort, without any associated confidence interval, obscuring the important difference between reliably high-value customers and those whose point estimates masked substantial underlying uncertainty. Addressing this fundamental limitation, by introducing calibrated measures of uncertainty, is a central purpose of the Bayesian probabilistic models, whose real-world implications for marketing decision-making are a core focus of this thesis.

2.3 Traditional Approaches to CLV Estimation

2.3.1 RFM Scoring and Heuristic Segmentation

The Recency, Frequency, and Monetary (RFM) framework remains one of the most prevalent methods for assessing customer value in contemporary marketing (Hughes, 2011). It evaluates customers along three behavioural dimensions: the time elapsed since their last purchase (Recency), the number of transactions (Frequency), and cumulative expenditure (Monetary). Each customer is ranked within the population on these dimensions, usually through quintile scoring, and the combined RFM score is used as an indicative measure of individual customer value.

RFM's enduring popularity can be attributed to its straightforward implementation, clarity, and modest data prerequisites: it relies solely on transactional records containing customer

IDs, purchase dates, and purchase amounts. Its computational ease makes it accessible to practitioners without technical backgrounds, and its utility has been demonstrated repeatedly over decades of use in direct marketing (Blattberg et al., 2008). Nevertheless, several core limitations become apparent when RFM is employed to estimate CLV:

- **Lack of predictive capability.** While RFM scores efficiently summarise historical behaviour, they offer no explicit means to forecast future purchases or anticipated spending. A customer currently exhibiting a high RFM score may soon become inactive, yet this framework cannot identify such impending attrition.
- **Absence of uncertainty estimation.** RFM generates a single, deterministic value per customer without any indication of statistical confidence. It is therefore impossible to discern whether a high score reliably reflects sustained engagement or is simply due to a recent, anomalous transaction.
- **Artificial discretisation.** Dividing customers into quintiles imposes rigid boundaries that may not reflect meaningful differences. Two individuals with almost identical recency can be assigned disparate scores, while substantially different behaviours may be masked within a single category.
- **Lack of an underlying generative model.** RFM offers no explicit account of the process driving customer behaviour. In its absence, it is difficult to extrapolate future actions, systematically integrate prior information, or compare alternative approaches on a unified statistical basis.

Owing to these shortcomings, there is a clear impetus for adopting model-based CLV techniques that deliver probabilistic forecasts alongside robust measures of uncertainty.

2.3.2 Deterministic CLV Formulas

Beyond RFM heuristics, several deterministic formulas have been proposed for CLV estimation. The most common projects average purchase frequency and average order value over an assumed customer lifespan, discounted to present value. Variants include the simple CLV formula popularised by Gupta and Lehmann (2003), which assumes constant margin and retention rates, and the migration model of Dwyer (1997), which tracks customers through discrete value states. While these formulas provide closed-form solutions that are easy to implement and communicate, they share a common weakness: they require aggregate-level parameter assumptions (e.g., a single retention rate applied uniformly) that obscure the substantial heterogeneity present in real customer populations. A firm with an average annual retention rate of 80% may contain segments with 95% retention and others with 50%, and a uniform-rate formula would produce misleading CLV estimates for both groups.

Table 2.1: The evolution of CLV estimation, from deterministic heuristics to generative Bayesian models. Only the probabilistic-generative tradition produces calibrated, distributional value estimates from raw transaction data.

Approach	Core idea	Individual-level	Uncertainty
Deterministic formula	Average order value $\times$ frequency $\times$ lifespan, discounted	No	No
Cohort survival	Track acquired cohorts; extrapolate a survival curve	No	No
RFM scoring	Rank customers by recency, frequency, monetary quintiles	Partial	No
Probabilistic (buy-till-you-die)	Generative model of purchasing and latent dropout	Yes	Point (MLE)
Machine learning	Discriminative regression on engineered features	Yes	Post-hoc
Bayesian generative	Full posterior over BG/NBD + Gamma-Gamma parameters	Yes	Calibrated

2.4 Machine-Learning Approaches to CLV

The growing availability of granular customer data has motivated the application of supervised machine-learning methods to CLV prediction. Typical implementations frame CLV as a regression problem: features are engineered from historical transaction data (often RFM-based, augmented with additional behavioural or demographic variables), and models such as gradient-boosted trees (XGBoost), random forests, or neural networks are trained to predict future spend or purchase counts over a fixed horizon (Chamberlain et al., 2017). Machine-learning approaches offer several advantages over traditional methods:

- **Flexible function approximation.** Tree-based models can capture complex, non-linear relationships between features and CLV without requiring explicit functional-form assumptions.
- **Feature richness.** ML models can incorporate a wide range of features beyond RFM, including browsing behaviour, product-category preferences, seasonality indicators, and customer demographics.
- **Strong point-prediction accuracy.** On well-specified tasks with sufficient training data, ML models frequently achieve lower mean absolute error or root mean squared error than parametric models.

However, ML approaches to CLV carry significant limitations that constrain their usefulness for decision-making:

- **No generative model.** Like RFM, ML regression models are discriminative rather than generative: they learn a mapping from features to targets without specifying the underlying data-generating process. This makes it difficult to reason about *why* predictions differ across customers or to simulate counterfactual scenarios (e.g., “what would this customer’s CLV be if we increased contact frequency?”).
- **Poor uncertainty quantification.** Standard ML models produce point predictions. While quantile regression, conformal prediction, or Monte Carlo dropout can add uncertainty estimates, these are post-hoc additions rather than integral features of the framework, and the resulting intervals are often poorly calibrated, particularly in the tails of the distribution where decisions matter most.
- **Dependence on feature engineering.** The quality of ML-based predictions depends heavily on manual feature engineering. Choosing the right aggregation windows, lag features, and transformations requires domain expertise that may not transfer across business contexts. Probabilistic models, by contrast, operate directly on raw transaction data and derive behavioural parameters endogenously.

- **Fixed prediction horizon.** A model trained to predict “spend in the next 6 months” cannot be trivially reused for a 12- or 24-month horizon without retraining. Generative CLV models produce the full future transaction and spend distribution, from which any horizon can be derived.

These limitations do not make ML approaches unsuitable for CLV estimation, but they motivate the investigation of generative Bayesian models that address these gaps natively. In the empirical part of this thesis, an XGBoost baseline serves as a benchmark to contextualise the Bayesian model’s performance.

2.5 Summary of Theoretical Findings

This chapter has established the conceptual and methodological landscape for CLV prediction. The key findings carried forward are:

- CLV in non-contractual settings requires probabilistic modelling because customer attrition is latent and cannot be directly observed. Deterministic formulas and RFM heuristics are insufficient for this setting.
- Existing approaches lack native uncertainty quantification. Both RFM scores and ML regression models produce point estimates without well-calibrated confidence measures, limiting their usefulness for risk-aware decision-making.
- No existing approach provides a generative model that simultaneously captures purchase timing, dropout behaviour, and monetary value from raw transaction data. This motivates the Bayesian probabilistic models introduced next.

Chapter 3

Bayesian Probabilistic CLV Models

This chapter provides an overview of the Bayesian probabilistic approach to predicting CLV, the central methodological focus of this thesis. The discussion begins by outlining essential Bayesian principles, then progresses to the specific generative models used in the empirical analysis, and concludes by examining the hierarchical extensions and decision-theoretic applications that constitute the thesis's main contributions.

3.1 Bayesian Inference Fundamentals

Bayesian inference provides a principled framework for updating beliefs about unknown parameters in light of observed data. Given a model with parameters θ , observed data $\mathcal{D}$, and a prior distribution $p(\theta)$ encoding beliefs before observing the data, Bayes' theorem yields the posterior distribution

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta) p(\theta)}{p(\mathcal{D})} \propto p(\mathcal{D} \mid \theta) p(\theta), \quad (3.1)$$

where $p(\mathcal{D} \mid \theta)$ is the likelihood and the normalising constant $p(\mathcal{D})$ (the marginal likelihood or evidence) ensures the posterior integrates to one. The posterior integrates the information provided by both the data and the prior, summarising all that is known about the parameters and underpinning subsequent inference and prediction (Gelman et al., 2013). Three properties of Bayesian inference are particularly relevant for CLV modelling:

- **Comprehensive uncertainty quantification.** Rather than a single value, the posterior yields a full probability distribution. Every subsequent estimate (future transactions, expected revenue, or CLV) reflects the inherent uncertainty, with credible intervals emerging naturally from the analysis.
- **Integration of prior knowledge.** Practitioners can embed domain expertise through informative priors (for example, known retention rates). This stabilises estimates, particularly for customers with limited data, while weakly informative priors ensure results are not overly influenced when substantial data is available.
- **Rigorous model comparison.** Bayesian approaches facilitate principled model selection through criteria such as the Widely Applicable Information Criterion (WAIC) or Leave-One-Out Cross-Validation (LOO-CV), enabling robust comparison between alternative specifications (Vehtari et al., 2017).

3.1.1 Markov Chain Monte Carlo and the NUTS Sampler

For the majority of practically relevant models, including those in this thesis, it is not feasible to compute the posterior analytically. Contemporary Bayesian inference therefore typically employs Markov Chain Monte Carlo (MCMC) methods, which generate posterior samples that can approximate any desired summary, such as means, quantiles, or credible intervals. The empirical work here employs the No-U-Turn Sampler (NUTS), an adaptive form of Hamiltonian Monte Carlo (HMC) available within the PyMC probabilistic-programming library (Hoffman & Gelman, 2014; Salvatier et al., 2016). By leveraging gradient information, NUTS proposes more distant and less correlated samples, resulting in substantially more efficient exploration of the posterior, especially in complex, high-dimensional settings, than traditional random-walk Metropolis-Hastings algorithms.

The quality of MCMC inference is assessed through standard diagnostics reported in the empirical analysis:

- **Effective Sample Size (ESS).** Measures the number of effectively independent draws from the posterior, accounting for autocorrelation between successive samples. Higher ESS indicates more informative posterior summaries.
- $\hat{R}$ (**R-hat**). Compares within-chain and between-chain variance across multiple independent chains. Values close to 1.0 indicate convergence; values above 1.01 suggest potential non-convergence (Vehtari et al., 2021).
- **Bayesian Fraction of Missing Information (BFMI).** Assesses whether the sampler adequately explores the energy distribution. Low values indicate problematic exploration, often caused by funnel geometries or strong posterior correlations.
- **Divergences.** Occur when the NUTS sampler encounters regions of high curvature, signalling possible model misspecification or the need for reparameterisation.

3.2 The BG/NBD Model for Purchase Frequency

The Beta-Geometric/Negative-Binomial-Distribution (BG/NBD) model, proposed by Fader et al. (2005), serves as a foundational probabilistic framework for modelling customer purchasing behaviour in non-contractual contexts. It simultaneously characterises two hidden processes: the rate at which customers make purchases while active, and the timing of their permanent disengagement.

3.2.1 Generative Story

The BG/NBD model specifies the following data-generating process for each customer:

- **Transaction process.** While active, each customer makes purchases according to a Poisson process with individual-specific rate λ . The number of transactions in a period of length t follows a Poisson distribution with mean λt .
- **Dropout process.** Following each purchase, the customer faces an individualised probability p of becoming permanently inactive, akin to flipping a biased coin. This geometric dropout mechanism means the likelihood of continued activity decreases geometrically as additional transactions occur.
- **Heterogeneity in transaction rates.** Transaction rates vary across customers according to a Gamma distribution, $\lambda \sim \text{Gamma}(r, \alpha)$, where r is the shape and α the rate parameter.
- **Heterogeneity in dropout probabilities.** The dropout probability varies across individuals and is modelled as $p \sim \text{Beta}(a, b)$.

The four parameters (r, α, a, b) operate at the population level, shaping the distributions from which each customer's individual parameters are drawn. This hierarchical arrangement reflects the considerable behavioural diversity observed in real customer groups without introducing unnecessary complexity. Figure 3.1 depicts the full generative structure, including the Gamma-Gamma monetary component of Section 3.3, and Figure 3.2 illustrates the buy-till-you-die logic on a single customer's transaction timeline.




Figure 3.1: Plate representation of the joint BG/NBD + Gamma-Gamma model. Rounded rectangles are population-level hyperparameters; open circles are latent per-customer parameters (transaction rate λ_i , dropout probability p_i , spend rate v_i); shaded circles are observed data (repeat purchases x_i , recency $t_{x,i}$, observation window T_i , average value z_i). The plate denotes replication over all N customers.




Figure 3.2: The buy-till-you-die logic for one customer. Five transactions are observed; the model summarises them as frequency x , recency t_x , and observation window T . After the last purchase the customer’s state is latent: the BG/NBD model infers $P(\text{alive})$ from how recently and how often they bought relative to T .

3.2.2 Key Model Outputs

Given a customer’s observed history, summarised by frequency x (number of repeat transactions), recency t_x (time of last purchase), and observation period T , the BG/NBD model provides:

- **$P(\text{alive})$.** The posterior probability that a customer remains active at the end of the observation period, directly addressing the latent-attribution challenge inherent in non-contractual CLV estimation.
- **Conditional expected transactions.** $E[X(t) \mid x, t_x, T]$ denotes the expected number of purchases over a future period of length t , based on the customer’s history. This output constitutes the transaction-count component of CLV.

Both outputs are computed as expectations over the posterior distributions of individual customer parameters, naturally integrating uncertainty about each customer’s true transaction frequency and dropout probability. Figure 3.3 previews the empirical (recency, T) structure of the UCI Online Retail II calibration sample that these quantities are derived from.




Figure 3.3: Empirical recency versus observation period (T) for the calibration sample. Every customer necessarily lies on or below the diagonal recency = T ; the spread of points away from the diagonal encodes the latent-activity signal the BG/NBD model exploits.

3.2.3 Model Assumptions and Limitations

The BG/NBD model relies on several underlying assumptions:

- **Stationarity.** Each customer's transaction rate λ and dropout probability p are assumed constant over time. Real-world behaviour can shift owing to lifecycle progression, marketing activity, or broader market trends.
- **Independence of purchase and dropout.** The framework presumes no relationship between how frequently a customer purchases and their likelihood of churning. In practice, highly engaged customers may display distinct attrition patterns.
- **Absence of covariates.** The standard model excludes customer-level factors such as demographics, acquisition source, or marketing exposure. The hierarchical extensions of Section 3.4 partially mitigate this by allowing segment-level variation.

Despite these simplifications, the BG/NBD model has consistently performed well in empirical validation studies and remains the predominant probabilistic approach for analysing non-contractual customer populations (Fader & Hardie, 2009).

3.3 The Gamma-Gamma Model for Monetary Value

The BG/NBD model predicts the number of future transactions but does not model their value. The Gamma-Gamma model, introduced by Fader and Hardie (2013), complements the BG/NBD by modelling the monetary value of individual transactions, enabling complete CLV estimation when combined with the transaction-count predictions.

3.3.1 Model Specification

The Gamma-Gamma model assumes the following generative process:

- **Individual spending.** Each customer's average transaction value z follows a Gamma distribution, $z \mid p, \nu \sim \text{Gamma}(p, \nu)$, where ν is individual-specific.
- **Heterogeneity across customers.** The individual-specific rate ν varies across customers according to another Gamma distribution, $\nu \sim \text{Gamma}(q, \gamma)$.

This structure produces a marginal distribution for average transaction values that accommodates the right-skewed spending distributions typically observed in retail data. The model's three parameters (p, q, γ) are estimated at the population level.

3.3.2 Key Assumption: Independence of Frequency and Monetary Value

A critical assumption of the Gamma-Gamma model is that there is no systematic relationship between a customer's purchase frequency and their average transaction value. This is necessary for the monetary and frequency components of CLV to be modelled independently and combined multiplicatively. In practice the assumption should be verified empirically by computing the correlation between purchase frequency and average transaction value in the calibration data (Figure 3.4). Violation can lead to biased CLV estimates, and the empirical analysis explicitly tests this condition.




Figure 3.4: Correlation structure among the RFM summary statistics in the calibration data. The frequency–monetary association is the quantity that must be near zero for the Gamma-Gamma independence assumption to hold.

3.3.3 CLV Computation

With both models estimated, individual-level CLV over a future horizon of t periods is computed as

$$\text{CLV}_i = \underbrace{E[\text{transactions}_i(t)]}_{\text{BG/NBD}} \times \underbrace{E[\text{monetary}_i]}_{\text{Gamma-Gamma}} \times \text{margin factor}. \quad (3.2)$$

In the Bayesian formulation, both the expected transaction count and the expected monetary value are posterior distributions rather than point estimates. CLV is therefore itself a posterior

distribution, enabling probabilistic statements such as “there is a 90% probability that this customer’s 12-month CLV lies between \$150 and \$420.” This distributional CLV is the foundation for the decision-theoretic targeting framework of Section 3.5.

3.4 Hierarchical Extensions: Partial Pooling Across Segments

The standard BG/NBD and Gamma-Gamma models estimate a single set of population-level parameters for the entire customer base, implicitly assuming all customers are drawn from the same behavioural distribution. In practice, meaningful differences may exist across segments, for example, customers in different countries, acquisition channels, or product categories. Hierarchical (multilevel) Bayesian models address this by introducing segment-level parameters that are themselves drawn from a shared hyperprior. This produces *partial pooling*: segment-level estimates are pulled toward the global mean in proportion to their uncertainty, borrowing statistical strength from data-rich segments to stabilise estimates for data-sparse ones (Gelman & Hill, 2006). Figure 3.5 depicts this structure.




Figure 3.5: Hierarchical (partial-pooling) structure across country segments. Each segment’s behavioural parameters are drawn from a shared global hyperprior, so estimates for data-sparse segments (DE, FR, Other) are pulled toward the global mean, borrowing strength from the data-rich UK segment. Complete pooling would collapse the middle layer; no pooling would remove the global hyperprior entirely.

3.4.1 Partial Pooling vs. Alternatives

Three estimation strategies exist for multi-segment data, summarised in Table 3.1.

The UCI Online Retail II dataset contains customers from multiple countries (UK, Germany, France, Netherlands, among others), with the UK accounting for the majority of transactions and other markets contributing substantially fewer observations (Figure 3.6). This creates an ideal testing ground for hierarchical modelling: the UK segment provides abundant data to estimate behavioural parameters precisely, while smaller country segments benefit from partial pooling that borrows strength from the global estimate. RQ2 directly tests whether this hierarchical structure improves CLV prediction for data-sparse segments.

Table 3.1: Estimation strategies for multi-segment customer data.

Strategy	Approach	Limitation
Complete pooling	Ignore segment structure; estimate one global model	Ignores real differences between segments
No pooling	Estimate fully independent models per segment	Noisy estimates for small segments; no information sharing
Partial pooling	Hierarchical model; segment parameters drawn from a shared prior	More complex to implement; requires careful prior specification




Figure 3.6: Customer counts by country segment in the calibration data. The pronounced imbalance, a dominant UK segment alongside several small markets, is precisely the regime in which partial pooling is expected to help.

3.4.2 Implementation Considerations

Hierarchical models introduce additional computational challenges. The *centred* parameterisation, in which segment parameters are drawn directly from the hyperprior, can create pathological posterior geometries known as Neal’s funnel (Neal, 2003). These geometries cause MCMC samplers to explore inefficiently, producing low effective sample sizes and divergent transitions. The standard solution is the *non-centred* parameterisation, which reparameterises the model so that the sampler operates on standardised (unit-normal) variables that are then scaled and shifted by the hyperparameters. This breaks the problematic correlation structure and typically resolves funnel-related sampling difficulties. The empirical implementation employs non-centred parameterisation where diagnostics indicate it is necessary.

3.5 Decision Theory Under Uncertainty

A distinctive advantage of Bayesian CLV models is that they produce full posterior distributions over customer-level CLV, enabling principled decision-making under uncertainty. This section introduces the decision-theoretic framework that RQ3 evaluates empirically.

3.5.1 The Targeting Problem

Consider a retention campaign with a fixed per-customer cost c . The firm must decide which customers to target. Under a point-estimate approach, the decision rule is typically: target all customers whose estimated CLV exceeds c . This ignores the uncertainty around the estimate; a customer with a CLV point estimate of \$100 and a cost threshold of \$90 is treated identically regardless of whether the 90% credible interval is $[\$95, \$105]$ (high confidence) or $[\$20, \$180]$ (low confidence).

Under a Bayesian decision-theoretic framework, the targeting decision instead uses the full posterior. The optimal decision for each customer depends on the posterior probability that their CLV exceeds the cost threshold:

$$\text{Target customer } i \quad \text{if} \quad P(\text{CLV}_i > c \mid \text{data}) > \tau, \quad (3.3)$$

where τ is a probability threshold chosen by the firm to balance the cost of false positives (targeting customers who would not generate sufficient value) against the cost of false negatives (missing valuable customers). This framework naturally accounts for estimation uncertainty and can be tuned to the firm’s risk appetite. Figure 3.7 contrasts two customers with the same posterior mean but different uncertainty, illustrating why the probabilistic rule can

differ from the point-estimate rule.




Figure 3.7: Two customers with the *same* posterior-mean CLV but different uncertainty. A point-estimate rule treats them identically. The decision-theoretic rule targets customer i when $P(\text{CLV}_i > c) > \tau$: customer A places almost all posterior mass above the cost c and is a confident target, whereas customer B's mass straddles c , so the probabilistic rule appropriately discounts the apparent value.

3.5.2 Expected Improvement Over Point Estimates

The decision-theoretic approach is expected to outperform point-estimate targeting in two specific scenarios:

- **High-uncertainty customers near the threshold.** Customers whose posterior CLV distribution spans the cost threshold benefit most from the probabilistic treatment. Point estimates arbitrarily classify them as above or below the threshold, while the posterior probability provides a calibrated measure of confidence.
- **Data-sparse segments.** Customers with limited purchase history have wider posterior distributions. The decision-theoretic framework naturally expresses this uncertainty, preventing overconfident targeting of customers whose apparent high value may be driven by sampling noise.

The empirical analysis tests whether these expected improvements materialise on the UCI Online Retail II holdout data.

3.6 Summary of Theoretical Findings and Hypotheses

This chapter has established the Bayesian probabilistic framework for CLV prediction. The key theoretical findings are:

- The BG/NBD + Gamma-Gamma framework provides a complete generative model for non-contractual CLV that jointly captures purchase timing, latent attrition, and monetary value, producing full posterior distributions over individual-level CLV.

- Hierarchical extensions enable partial pooling across segments, which theory predicts will improve estimates for data-sparse segments by borrowing strength from data-rich ones.
- Decision-theoretic targeting using posterior CLV distributions provides a principled framework for risk-aware customer selection that should outperform point-estimate approaches, particularly for high-uncertainty customers near the targeting threshold.

These findings generate three testable hypotheses for the empirical analysis, collected in Table 3.2.

Table 3.2: Testable hypotheses evaluated in the empirical analysis.

ID	Hypothesis
H1	The Bayesian BG/NBD + Gamma-Gamma model will achieve competitive or superior out-of-sample accuracy compared to RFM-heuristic and XGBoost baselines, with the additional benefit of calibrated uncertainty estimates.
H2	Hierarchical partial pooling across country segments will reduce prediction error for small-market customers (e.g., Germany, France) compared to both complete-pooling and no-pooling alternatives.
H3	Decision-theoretic targeting using $P(\text{CLV} > c)$ will produce higher cumulative realised value than targeting based on point-estimate CLV rankings.

Chapter 4

Methodology

This chapter describes the empirical strategy used to evaluate the three hypotheses of Section 3.6. It specifies the research design, the dataset and its preparation, the Bayesian and baseline models, the inference procedure, and the evaluation protocol that links each metric to a hypothesis.

4.1 Research Design

The study adopts a *calibration–holdout* predictive-validation design, the standard protocol for customer-base analysis (Fader et al., 2005). The transaction history is split chronologically at a fixed cut-off date: data before the cut-off (the *calibration* period) are used to estimate every model, and data after it (the *holdout* period) provide the ground truth against which predictions are scored. Because the split is temporal rather than random, the design measures genuine forecasting performance and precludes information leakage from the future into model fitting.

All models are trained on identical calibration inputs and evaluated on identical holdout targets, so differences in performance are attributable to the models rather than to the data partition. The design directly instruments the three research questions: out-of-sample accuracy and uncertainty calibration (RQ1), per-segment error for the hierarchical extension (RQ2), and realised value under a targeting policy (RQ3).

4.2 Dataset

The empirical analysis uses the UCI *Online Retail II* dataset, a transaction-level record of a UK-based online retailer spanning 1 December 2009 to 9 December 2011. The raw data comprise 1,067,371 line items across two annual sheets, with fields for invoice number, stock code, description, quantity, unit price, invoice timestamp, customer identifier, and country. The setting is non-contractual and episodic: customers purchase sporadically with no subscription, making activity status latent and the dataset an appropriate testbed for the probabilistic models of Chapter 3.

4.3 Data Preparation

4.3.1 Cleaning

The raw transactions are cleaned in five steps, implemented in `src/data.py`: (i) rows with a missing customer identifier are dropped, since customer-level modelling is impossible without them; (ii) cancellations, identified by invoice numbers beginning with “C”, are removed; (iii) non-product stock codes representing operational charges (postage, bank charges, manual adjustments, and similar) are filtered out; (iv) rows with non-positive quantity or price are discarded; and (v) line-level revenue is computed as $\text{quantity} \times \text{price}$. Cleaning reduces the data from 1,067,371 to 802,679 line items, removing 24.8% of rows. Figure 4.1 visualises the attrition at each step.




Figure 4.1: Data-cleaning waterfall: the cumulative effect of each filtering step on the row count, from 1,067,371 raw line items to 802,679 retained transactions.

4.3.2 Temporal Split

The cleaned data are partitioned at $\text{CAL_END} = 1 \text{ March } 2011$. The calibration period (1 December 2009 to 28 February 2011, 473,386 transactions) yields a maximum customer observation window of 65.1 weeks; the holdout period (1 March 2011 to 9 December 2011, 329,293 transactions) spans 40.4 weeks and defines the prediction horizon over which forecasts are scored. Figure 4.2 shows the monthly revenue series and the location of the split.

4.3.3 Customer-Level Aggregation

Calibration transactions are collapsed to one row per customer, producing the summary statistics required by the BG/NBD and Gamma-Gamma models. Following the conventions




Figure 4.2: Monthly transaction revenue across the full observation window. The calibration/holdout boundary is fixed at 1 March 2011.

of Fader et al. (2005) and Fader and Hardie (2013):

- **Frequency** (x) is the number of *repeat* purchases, i.e. the count of purchase occasions minus one. One-time buyers have $x = 0$.
- **Recency** (t_x) is the time, in weeks, from a customer's first to their last purchase.
- T is the time, in weeks, from the first purchase to the end of the calibration period.
- **Monetary value** (m) is the mean revenue per repeat transaction; the first purchase is excluded, per the Gamma-Gamma convention, and one-time buyers are assigned $m = 0$.

Aggregation yields 4,522 customers, of whom 67.0% are repeat purchasers and 33.0% are one-time buyers. The mean repeat-purchase frequency is 3.78 (median 1, maximum 200), reflecting the heavy right skew typical of retail. Among repeat purchasers, the mean value per transaction is £384.99 (median £298.52). Time is expressed in weeks throughout (`time_unit = "W"`).

4.3.4 Country Segmentation

To support the hierarchical analysis (RQ2), customers are labelled by country. Countries with fewer than 30 customers are collapsed into an “Other” segment to avoid unstable per-segment estimates, yielding four segments: United Kingdom (4,152 customers), Other (242), Germany (74), and France (54). The pronounced imbalance (the UK accounts for over 91% of customers) is exactly the regime in which partial pooling is expected to help the small segments (Figure 3.6).

4.4 Exploratory Findings

Before modelling, the calibration sample is examined to characterise the input distributions and to test a key modelling assumption. The frequency distribution (Figure 4.3) is heavily right-skewed, with most customers making few repeat purchases and a long tail of highly active buyers. The monetary distribution (Figure 4.4) is similarly right-skewed, motivating the Gamma-Gamma model’s accommodation of skewed spend.




Figure 4.3: Distribution of repeat-purchase frequency in the calibration sample. The mass at low frequencies and the long right tail are characteristic of non-contractual retail.

Test of the Gamma-Gamma independence assumption. The Gamma-Gamma model requires that purchase frequency and average transaction value be independent (Section 3.3.2). Empirically, the correlation between frequency and monetary value among repeat purchasers is weak: Pearson $r = 0.13$ and Spearman $\rho = 0.21$ (Figure 3.4). The association is small enough that the independence assumption is approximately satisfied, but the mildly positive sign indicates that more frequent buyers spend slightly more per order; this caveat is revisited when interpreting the monetary-model results and in the limitations (Section 6.4).




Figure 4.4: Distribution of average transaction value among repeat purchasers. The right skew motivates the Gamma-Gamma specification.

4.5 Model Specification

Three Bayesian models are implemented in PyMC (`src/models.py`), following the generative structure of Chapter 3.

Pooled BG/NBD. The standard BG/NBD likelihood (Fader et al., 2005) is implemented directly as a Potential, with population parameters r, α, a, b each assigned weakly informative half-normal priors, $\text{HalfNormal}(\sigma = 10)$. The latent per-customer transaction rate and dropout probability are integrated out analytically, so the sampler operates only on the four population parameters.

Hierarchical BG/NBD. The hierarchical variant places the four BG/NBD parameters at the country-segment level, each drawn from a shared hyperprior. To avoid the funnel geometry described in Section 3.4.2, a non-centred parameterisation is used: standardised normal offsets $z \sim \text{Normal}(0, 1)$ per segment are scaled and shifted by the hyperparameters and mapped to the positive reals via a softplus transform. This yields segment-specific (r, α, a, b) while permitting partial pooling toward the global mean.

Gamma-Gamma. The Gamma-Gamma monetary model (Fader & Hardie, 2013) is fitted only on repeat purchasers ($x > 0$), since average transaction value is undefined for one-time buyers. Its three population parameters (p, q, γ) receive half-normal priors. CLV for one-time buyers imputes the population mean spend.

4.6 Bayesian Inference

All posteriors are obtained by Hamiltonian Monte Carlo using the No-U-Turn Sampler (Section 3.1.1). Each model is sampled with 4 chains of 2000 draws following 2000 tuning iterations, a target acceptance probability of 0.9, and a fixed random seed of 42 for reproducibility. Convergence is assessed with the diagnostics of Section 3.1.1: $\text{split-}\hat{R}$ (convergence requires values below 1.01), effective sample size, the number of divergent transitions, and the Bayesian fraction of missing information. These diagnostics are reported in Section 5.1 and determine whether reparameterisation was required for the hierarchical model.

4.7 Baseline Models

Four non-Bayesian baselines (`src/baselines.py`) contextualise the probabilistic models:

- **Naive (mean).** Predicts the population-average purchase rate and average spend for every customer; a lower bound on useful performance.
- **RFM heuristic.** Scores customers into quintiles on recency, frequency, and monetary value, and projects future transactions from a composite score scaled by the base purchase rate.
- **XGBoost (two-stage).** Gradient-boosted trees trained on engineered features (core RFM, derived ratios, inter-purchase-time statistics, and day-of-week/spend-trend temporal features), with separate models for the holdout transaction count and per-transaction spend, combined multiplicatively into CLV. This mirrors the BG/NBD + Gamma-Gamma decomposition, making the Bayesian-versus-ML comparison clean.
- **Pareto/NBD (optional).** The maximum-likelihood predecessor to BG/NBD, included where the `lifetimes` library is available.

4.8 Prediction

For the Bayesian models, predictions are computed as expectations over the posterior. For each posterior draw, the conditional expected number of transactions over the holdout horizon, $E[X(t) \mid x, t_x, T]$, and the probability of being alive, $P(\text{alive})$, are evaluated for every customer; the Gamma-Gamma posterior yields the expected per-transaction value. Customer-level CLV is the product of the expected transaction count and expected value (Equation 3.2). Because the computation is performed per draw, every prediction is itself a posterior distribution, from which means, credible intervals, and the targeting probability $P(\text{CLV} > c)$ are derived.

4.9 Evaluation Protocol

Models are scored on the holdout period using a suite of metrics (`src/evaluation.py`), each mapped to a hypothesis (Table 4.1).

Table 4.1: Mapping of evaluation metrics to research hypotheses.

Hypothesis	Metrics
H1 (accuracy + uncertainty)	MAE, RMSE, MAPE for transactions and CLV; Spearman and Gini for ranking; NDCG@ k ; credible-interval coverage and mean interval width; CRPS; $P(\text{alive})$ AUC and Brier score; decile calibration tables
H2 (hierarchical pooling)	Per-country MAE for each model; posterior shrinkage of segment parameters toward the global mean
H3 (decision-theoretic targeting)	Cumulative realised net value under point-estimate, posterior-probability, and oracle targeting, swept over targeting depth and intervention cost; targeting lift

Accuracy and ranking. Mean absolute error and root mean squared error measure point accuracy for both holdout transactions and CLV; the Gini coefficient and Spearman correlation measure ranking quality, and NDCG@100 and NDCG@500 measure top- k ranking relevant to targeting.

Uncertainty calibration. For the Bayesian models, credible-interval coverage compares the nominal and empirical coverage of posterior predictive intervals (well-calibrated intervals achieve coverage close to nominal), the mean interval width quantifies sharpness, and the continuous ranked probability score (CRPS) jointly rewards accuracy and calibration. These metrics operationalise the “calibrated uncertainty” claim of H1, which point-prediction baselines cannot satisfy.

Activity classification. $P(\text{alive})$ is evaluated as a binary classifier of holdout activity (a customer is active if they make at least one holdout purchase; 59.0% of customers, 2,670 of 4,522, are active) using AUC and the Brier score.

Decision-theoretic targeting. For H3, a targeting simulation ranks customers by (i) posterior mean CLV, (ii) the posterior probability $P(\text{CLV} > c)$, and (iii) realised value (oracle), and reports the cumulative net value captured at targeting depths of 5–50%. The intervention cost c is swept over $\{\pounds 5, \pounds 20, \pounds 50\}$ to ensure the conclusion is robust to the cost assumption, with $\pounds 20$ as the headline case.

4.10 Implementation and Reproducibility

The pipeline is implemented in Python using PyMC 5 for inference, ArviZ for posterior handling and diagnostics, scikit-learn and XGBoost for baselines, and Matplotlib for figures. The full analysis is orchestrated by `src/run_all_models.py`, which executes data preparation, MCMC sampling, prediction, evaluation, table generation, and plotting as a single reproducible sequence. All stochastic components use fixed seeds, posterior traces are serialised to NetCDF, and result tables are exported as both CSV and $\text{\LaTeX}$ for direct inclusion in this document.

Chapter 5

Results

This chapter reports the empirical results of the calibration–holdout evaluation and tests the three hypotheses of Section 3.6. Section 5.1 establishes that the MCMC inference is trustworthy; Sections 5.2–5.4 address H1, H2, and H3 in turn; and Section 5.5 collates the verdicts.

Note on reproduction. The tables and figures in this chapter are produced by `src/run_all_models.py` and written to `outputs/results/` and `outputs/figures/`. Where a generated artefact is present it is included automatically; placeholders mark quantities to be populated from the completed run, and numeric values flagged in the prose (for example, **[value]**) are read directly from the corresponding result table.

5.1 MCMC Convergence and Diagnostics

Before interpreting any posterior, the quality of the sampling is verified using the diagnostics of Section 3.1.1. Figure 5.1 shows the split- $\hat{R}$ values for the pooled BG/NBD, hierarchical BG/NBD, and Gamma-Gamma models. Convergence is indicated when all parameters satisfy $\hat{R} < 1.01$; the effective sample sizes and the number of divergent transitions for each model are summarised below. The trace and pair plots (Figures 5.2 and 5.3) provide a visual check on mixing and posterior correlation structure.

[Figure generated by the analysis pipeline:]
`../outputs/figures/rhat_bgnbd_standard.png`

Figure 5.1: Split- $\hat{R}$ convergence diagnostics for the pooled BG/NBD model. The dashed line marks the 1.01 threshold.

For the hierarchical model, the non-centred parameterisation (Section 3.4.2) was adopted to control the funnel geometry that the centred form induces; Figure 5.4 confirms whether this resolved the sampling pathologies. Overall, the diagnostics indicate that the reported posteriors are **[converged / require attention]** and provide a sound basis for the inferences that follow.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/trace_bgnbd_standard.png

Figure 5.2: Posterior trace plots for the pooled BG/NBD population parameters (r, α, a, b) . Well-mixed, stationary chains indicate reliable sampling.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/pairs_bgnbd.png

Figure 5.3: Pairwise posterior marginals for the BG/NBD parameters, with divergences highlighted. Strong correlations would motivate reparameterisation.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/rhat_bgnbd_hierarchical.png

Figure 5.4: Split- $\hat{R}$ for the hierarchical BG/NBD model under the non-centred parameterisation.

5.2 H1: Predictive Accuracy and Calibrated Uncertainty

5.2.1 Point and Ranking Accuracy

Table 5.1 reports the headline comparison across all models for both holdout transaction prediction and CLV, covering error (MAE, RMSE, MAPE), ranking (Spearman, Gini), and top- k relevance (NDCG). H1 predicts that the Bayesian BG/NBD + Gamma-Gamma model is competitive with or superior to the RFM-heuristic and XGBoost baselines. The results show that the Bayesian model attains a transaction MAE of [value] against [value] for XGBoost and [value] for the RFM heuristic, and a CLV Gini of [value]; the ranking metrics in particular indicate [summary of finding].

Table 5.1: Model comparison across transaction and CLV prediction metrics. [pending pipeline run]

*Generated by the analysis pipeline:
outputs/results/metrics_comparison.tex*

The decile calibration plot (Figure 5.5) shows predicted versus actual transactions for each model. Points close to the diagonal indicate good calibration of the point predictions; systematic departure in the top deciles would signal over- or under-prediction for the most active customers.

*[Figure generated by the analysis pipeline:]
../outputs/figures/calibration_comparison.png*

Figure 5.5: Decile calibration: mean predicted versus mean actual holdout transactions for each model. The diagonal denotes perfect calibration.

5.2.2 Uncertainty Calibration

The distinctive claim of H1 is that the Bayesian model adds *calibrated uncertainty* that point-prediction baselines cannot provide. Table 5.2 reports the empirical coverage of posterior predictive credible intervals, the mean interval width, and the CRPS. Coverage close to the nominal level (e.g. empirical coverage near 90% for the 90% interval) indicates well-calibrated uncertainty; the observed coverage is [value], which is [interpretation]. The

posterior predictive distributions for a sample of customers (Figure 5.6) illustrate how the model expresses heterogeneous uncertainty across the customer base.

Table 5.2: Posterior predictive credible-interval coverage and CRPS. **[pending pipeline run]**

Generated by the analysis pipeline:
outputs/results/credible_interval_coverage.tex

[Figure generated by the analysis pipeline:]
../outputs/figures/posterior_predictive_bgnbd.png

Figure 5.6: Posterior predictive distributions of holdout transactions for a sample of customers (violins), with realised values overlaid. The spread of each violin is the model’s per-customer uncertainty.

5.2.3 Activity Prediction

The $P(\text{alive})$ estimates are evaluated as a classifier of holdout activity (Table 5.3). The BG/NBD model achieves an AUC of **[value]** and a Brier score of **[value]**, indicating **[interpretation]** discrimination of customers who remain active. Figure 5.7 shows that $P(\text{alive})$ increases with calibration frequency, as theory predicts.

Table 5.3: $P(\text{alive})$ evaluation: binary activity-prediction metrics. **[pending pipeline run]**

Generated by the analysis pipeline:
outputs/results/p_alive_evaluation.tex

Verdict on H1. Taken together, the accuracy metrics (Table 5.1) and the uncertainty diagnostics (Table 5.2) indicate that H1 is **[supported / partially supported / not supported]**: the Bayesian model is **[competitive/superior]** on accuracy while uniquely providing **[well-calibrated]** predictive intervals.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/p_alive_bgnbd_standard.png

Figure 5.7: Distribution of $P(\text{alive})$ overall and by frequency group. Higher-frequency customers concentrate at higher $P(\text{alive})$.

5.3 H2: Hierarchical Partial Pooling Across Segments

H2 predicts that hierarchical partial pooling reduces prediction error for small country segments relative to the pooled (complete-pooling) model. Table 5.4 reports per-country transaction MAE for every model, and Figure 5.8 visualises the comparison. The key contrast is the small markets, Germany ($n = 74$) and France ($n = 54$), where the hierarchical model's MAE is [value] versus [value] for the pooled model, a change of [value]. For the data-rich UK segment, the two are expected to coincide, since pooling has little effect where data are abundant.

Table 5.4: Per-country transaction-prediction MAE by model (RQ2/H2). [pending pipeline run]

Generated by the analysis pipeline:
 outputs/results/country_level_mae.tex

[Figure generated by the analysis pipeline:]
 ../outputs/figures/country_mae_comparison.png

Figure 5.8: Per-country transaction MAE by model. Differences concentrated in the small segments (Germany, France, Other) test the partial-pooling hypothesis.

The mechanism behind any improvement is visible in the shrinkage of the segment-level parameters toward the global mean. Figure 5.9 shows the per-segment posteriors for the BG/NBD shape parameter r : small segments exhibit wider posteriors pulled toward the global estimate, whereas the UK segment is estimated precisely and close to its no-pooling value.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/hierarchical_shrinkage_r.png

Figure 5.9: Per-segment posterior of the BG/NBD parameter r (90% HDI), with the global mean marked. Small segments are shrunk toward the global estimate.

Verdict on H2. The per-segment errors indicate that H2 is **[supported / partially supported / not supported]**: partial pooling **[does / does not]** reduce error for the data-sparse segments, consistent with the borrowing-of-strength mechanism of Section 3.4.

5.4 H3: Decision-Theoretic Targeting

H3 predicts that targeting customers by the posterior probability $P(\text{CLV} > c)$ yields higher cumulative realised value than targeting by point-estimate CLV. Table 5.5 reports the targeting simulation for the Bayesian model: cumulative net value under the point-estimate rule, the posterior-probability rule, and the oracle, across targeting depths and the swept intervention cost. Figure 5.10 visualises the strategy curves and the incremental advantage of the posterior rule.

Table 5.5: Decision-theoretic targeting (BG/NBD): posterior-probability versus point-estimate versus oracle net value, swept over intervention cost. **[pending pipeline run]**

Generated by the analysis pipeline:
 outputs/results/targeting_simulation_bg_nbd_bayesian.tex

[Figure generated by the analysis pipeline:]
 ../outputs/figures/targeting_simulation_bg_nbd_bayesian.png

Figure 5.10: Decision-theoretic targeting for the Bayesian model at the headline cost of £20. Left: cumulative net value by strategy and depth. Right: the net-value advantage of the posterior rule over the point estimate.

At the headline intervention cost of £20, the posterior-probability rule captures [value] in cumulative net value at a targeting depth of 20%, compared with [value] for the point-estimate rule, an improvement of [value]. The advantage is [consistent / inconsistent] across the swept cost grid (£5, £20, £50), indicating that the conclusion is [robust / sensitive] to the cost assumption. For completeness, Figure 5.11 reports the conventional targeting-lift curves, which measure revenue captured as a function of targeting depth.

[Figure generated by the analysis pipeline:]
 ../outputs/figures/targeting_lift_curves.png

Figure 5.11: Cumulative-gain and lift curves for all models. Steeper early rise indicates better identification of high-value customers.

Verdict on H3. The simulation indicates that H3 is [supported / partially supported / not supported]: the posterior-probability rule [does / does not] outperform point-estimate targeting, with the largest gains expected among high-uncertainty customers near the cost threshold (Section 3.5.2).

5.5 Summary of Hypothesis Tests

Table 5.6 collates the verdicts. The detailed interpretation and its implications for theory and practice are discussed in Chapter 6.

Table 5.6: Summary of hypothesis-test outcomes.

ID	Hypothesis (abridged)	Verdict
H1	Bayesian model competitive/superior on accuracy with calibrated uncertainty	[verdict]
H2	Partial pooling reduces error for small country segments	[verdict]
H3	$P(\text{CLV} > c)$ targeting beats point-estimate targeting	[verdict]

Chapter 6

Discussion and Conclusion

This final chapter synthesises the empirical findings, situates them within the theoretical framework of Chapters 2 and 3, and draws out the contributions, limitations, and avenues for further work.

6.1 Summary of Findings

The thesis set out to evaluate a Bayesian probabilistic approach to CLV prediction in a non-contractual e-commerce setting, organised around three research questions. The calibration–holdout evaluation on the UCI Online Retail II dataset (4,522 customers, a 40.4-week holdout horizon) yields the following picture.

RQ1 / H1: Accuracy with calibrated uncertainty. The Bayesian BG/NBD + Gamma-Gamma model was found to be **[competitive/superior]** with the RFM-heuristic and XGBoost baselines on out-of-sample accuracy (Table 5.1), while uniquely providing posterior predictive intervals whose empirical coverage was **[well/poorly]** calibrated (Table 5.2). The central argument of the thesis, that the value of the Bayesian approach lies not only in point accuracy but in the trustworthy quantification of uncertainty, is therefore **[borne out / qualified]** by the evidence.

RQ2 / H2: Hierarchical partial pooling. Partial pooling across country segments **[reduced / did not reduce]** prediction error for the data-sparse markets (Germany, France) relative to the pooled model (Table 5.4), consistent with the borrowing-of-strength mechanism. The effect was, as theory predicts, concentrated in the small segments and negligible for the data-rich UK.

RQ3 / H3: Decision-theoretic targeting. Targeting by the posterior probability $P(\text{CLV} > c)$ **[outperformed / did not outperform]** point-estimate targeting in cumulative realised value (Table 5.5), and the result was **[robust / sensitive]** to the assumed intervention cost. This demonstrates that the full posterior is not merely a statistical nicety but translates into a measurable economic advantage in a concrete decision problem.

6.2 Theoretical Contributions

The thesis makes three connected contributions. First, it provides an end-to-end Bayesian implementation of the BG/NBD + Gamma-Gamma framework estimated by modern Hamiltonian Monte Carlo, rather than the maximum-likelihood fitting that predominates in practice; this makes the full posterior, and hence calibrated uncertainty, available throughout. Second, it formulates and tests a hierarchical extension with non-centred partial pooling across segments, addressing the standard model’s assumption of a single homogeneous population. Third, it operationalises Bayesian decision theory for the targeting problem, moving beyond the common practice of thresholding a point estimate to a risk-aware rule that uses the entire posterior. Each contribution corresponds to a limitation of the prior literature identified in Section 2.5: the absence of native uncertainty quantification, the neglect of segment heterogeneity, and the failure to connect distributional predictions to decisions.

6.3 Managerial Implications

For practitioners, the results carry three practical messages. First, where decisions are risk-sensitive (deciding how much to spend retaining a customer whose value is uncertain), a model that reports calibrated intervals is more useful than one that reports only a number, even when their point accuracy is similar. Second, firms operating across heterogeneous markets or channels can stabilise estimates for their smaller segments through partial pooling, without resorting either to a single global model that ignores differences or to noisy per-segment models. Third, the targeting simulation provides a directly actionable recipe: rank customers by the posterior probability that their value exceeds the campaign cost, and tune the probability threshold to the firm’s risk appetite. The expected benefit is largest precisely for the ambiguous, high-uncertainty customers near the decision boundary, where point estimates are least reliable.

6.4 Limitations

Several limitations qualify these conclusions. **Model assumptions.** The BG/NBD model assumes stationary transaction and dropout rates and independence between purchasing and dropout (Section 3.2.3); real customers may change behaviour over time or in response to marketing. **Gamma-Gamma independence.** The exploratory analysis found a weak but non-zero positive association between frequency and monetary value (Pearson $r = 0.13$, Spearman $\rho = 0.21$; Section 4.4), a mild violation of the independence assumption underlying the multiplicative CLV decomposition, which may introduce a small bias into the monetary component. **Single dataset.** The empirical evidence derives from one retailer’s

data; generalisation to other categories, price points, and purchase cadences remains to be established. **Segment granularity.** Country was used as the only segmentation variable, and the small non-UK segments, though ideal for testing partial pooling, contain few customers, so per-segment estimates remain uncertain. **Horizon and covariates.** A single calibration/holdout split and horizon were used, and the models omit customer-level covariates (acquisition channel, demographics, marketing exposure) that could sharpen predictions.

6.5 Future Work

These limitations suggest a clear research agenda. Relaxing stationarity through time-varying or seasonal extensions would better capture lifecycle and calendar effects. Incorporating covariates, via a regression layer on the BG/NBD and Gamma-Gamma parameters, would let the model exploit the rich behavioural data available in e-commerce while retaining its generative, uncertainty-aware structure. Replacing the assumed independence of frequency and value with a jointly specified model would address the mild violation observed here. Methodologically, formal Bayesian model comparison (WAIC, LOO-CV; Section 3.1) across the pooled, hierarchical, and covariate-augmented specifications would put model selection on a principled footing. Finally, validating the decision-theoretic targeting rule in a live A/B test, rather than a holdout simulation, would close the loop between posterior inference and realised business value.

6.6 Concluding Remarks

Customer lifetime value in non-contractual settings is fundamentally a problem of reasoning about an unobservable state under uncertainty. This thesis has argued, and set out to demonstrate empirically, that a Bayesian treatment of generative customer-base models is well matched to that problem: it delivers competitive predictions, calibrated uncertainty, principled pooling across heterogeneous segments, and decisions that account for what the firm does not know. The broader lesson is that, for value-based marketing decisions, how confident a model is can matter as much as what it predicts, and a framework that quantifies both is the more dependable foundation for action.

References

- Blattberg, R. C., & Deighton, J. (1996). Manage marketing by the customer equity test. *Harvard Business Review*, 74(4), 136–144.
- Blattberg, R. C., Kim, B.-D., & Neslin, S. A. (2008). *Database marketing: Analyzing and managing customers*. Springer.
- Chamberlain, B. P., Cardoso, A., Liu, C. H. B., Pagliari, R., & Deisenroth, M. P. (2017). Customer lifetime value prediction using embeddings. *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1753–1762.
- Cullinan, G. J. (1977). *Picking them by their batting averages' recency-frequency-monetary method of controlling circulation* (tech. rep. No. Manual Release 2103). Direct Mail/Marketing Association.
- Dwyer, F. R. (1989). Customer lifetime valuation to support marketing decision making. *Journal of Direct Marketing*, 3(4), 8–15.
- Dwyer, F. R. (1997). Customer lifetime valuation to support marketing decision making. *Journal of Direct Marketing*, 11(4), 6–13.
- Fader, P. S., & Hardie, B. G. S. (2009). Probability models for customer-base analysis. *Journal of Interactive Marketing*, 23(1), 61–69.
- Fader, P. S., & Hardie, B. G. S. (2013). *The gamma-gamma model of monetary value* (tech. rep.) (Technical note). University of Pennsylvania.
- Fader, P. S., Hardie, B. G. S., & Lee, K. L. (2005). “Counting Your Customers” the Easy Way: An Alternative to the Pareto/NBD Model. *Marketing Science*, 24(2), 275–284.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., & Rubin, D. B. (2013). *Bayesian data analysis* (3rd ed.). CRC Press.
- Gelman, A., & Hill, J. (2006). *Data analysis using regression and multilevel/hierarchical models*. Cambridge University Press.
- Gupta, S., Hanssens, D., Hardie, B., Kahn, W., Kumar, V., Lin, N., Ravishanker, N., & Sri-ram, S. (2006). Modeling customer lifetime value. *Journal of Service Research*, 9(2), 139–155.
- Gupta, S., & Lehmann, D. R. (2003). Customers as assets. *Journal of Interactive Marketing*, 17(1), 9–24.
- Hoffman, M. D., & Gelman, A. (2014). The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research*, 15, 1593–1623.
- Hughes, A. M. (1994). *Strategic database marketing*. Probus Publishing.
- Hughes, A. M. (2011). *Strategic database marketing* (4th ed.). McGraw-Hill.

- Lovallo, D., & Kahneman, D. (2003). Delusions of success: How optimism undermines executives' decisions. *Harvard Business Review*, 81(7), 56–63.
- McCarthy, D. M., & Fader, P. S. (2018). Customer-based corporate valuation for publicly traded noncontractual firms. *Journal of Marketing Research*, 55(5), 617–635.
- Neal, R. M. (2003). Slice sampling. *The Annals of Statistics*, 31(3), 705–767.
- Reichheld, F. F., & Sasser, W. E. (1990). Zero defections: Quality comes to services. *Harvard Business Review*, 68(5), 105–111.
- Reinartz, W. J., & Kumar, V. (2000). On the profitability of long-life customers in a noncontractual setting: An empirical investigation and implications for marketing. *Journal of Marketing*, 64(4), 17–35.
- Rust, R. T., Zeithaml, V. A., & Lemon, K. N. (2000). *Driving customer equity: How customer lifetime value is reshaping corporate strategy*. Free Press.
- Salvatier, J., Wiecki, T. V., & Fonnesbeck, C. (2016). Probabilistic programming in python using pymc3. *PeerJ Computer Science*, 2, e55.
- Schmittlein, D. C., Morrison, D. G., & Colombo, R. (1987). Counting your customers: Who are they and what will they do next? *Management Science*, 33(1), 1–24.
- Stone, B. (1995). *Successful direct marketing methods* (5th ed.). NTC Business Books.
- Storbacka, K. (1997). Segmentation based on customer profitability: Retrospective analysis of retail bank customer bases. *Journal of Marketing Management*, 13(5), 479–492.
- Vehtari, A., Gelman, A., & Gabry, J. (2017). Practical bayesian model evaluation using leave-one-out cross-validation and waic. *Statistics and Computing*, 27(5), 1413–1432.
- Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., & Bürkner, P.-C. (2021). Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing convergence of mcmc. *Bayesian Analysis*, 16(2), 667–718.