
Customer Lifetime Value Prediction

Using Bayesian Methods

A Probabilistic Approach to Non-Contractual
Customer-Base Analysis in E-Commerce

Master's Thesis

submitted in partial fulfilment of the requirements
for the degree of Master of Science

Devansh

July 26, 2026

Abstract

Customer Lifetime Value (CLV) is the expected net present value of all future cash flows from a customer relationship and is a foundational input to marketing-resource allocation. In non-contractual settings such as e-commerce, customers never announce their departure, so the central modelling difficulty is probabilistic reasoning about an unobservable attrition state. This thesis develops and evaluates a Bayesian probabilistic approach to CLV prediction built on the Beta-Geometric/Negative-Binomial-Distribution (BG/NBD) model for purchase frequency and the Gamma-Gamma model for monetary value, estimated by Hamiltonian Monte Carlo with the No-U-Turn Sampler. Three contributions are examined empirically on the UCI Online Retail II dataset: (i) whether the Bayesian model matches or exceeds RFM and XGBoost baselines on out-of-sample accuracy while additionally providing calibrated uncertainty (RQ1/H1); (ii) whether a hierarchical extension with partial pooling across country segments improves prediction for data-sparse markets (RQ2/H2); and (iii) whether decision-theoretic targeting based on the posterior probability $P(\text{CLV} > c)$ realises greater value than targeting by point-estimate CLV (RQ3/H3). On a 40-week hold-out the Bayesian model outperforms all baselines on accuracy and ranking (transaction MAE 1.74 against 2.12 for XGBoost; NDCG@100 0.91 against 0.53) while providing posterior predictive intervals with 89.3% empirical coverage at the nominal 90% level, supporting H1. Partial pooling is never the worst of the three pooling strategies and achieves the lowest aggregated error across the small segments, but does not improve on complete pooling, so H2 is partially supported. Probability-based targeting matches but does not beat point-estimate targeting (H3 not supported); the analysis instead surfaces a decision-relevant methodological pitfall, namely that targeting probabilities must be computed from the posterior predictive distribution rather than from the posterior of expected CLV.

Contents

Abstract

List of Abbreviations viii

Notation ix

1 Introduction 1

1.1	Background and Motivation	1
1.2	Understanding Customer Lifetime Value: Concept and Historical Development	1
1.3	Customer Lifetime Value in E-Commerce and Modern Business	2
1.4	Research Questions and Hypotheses	3
1.5	Structure of the Thesis	3

2 Theoretical Foundations of Customer Lifetime Value 4

2.1	Customer Lifetime Value: Concept	4
2.1.1	Contractual vs. Non-Contractual Settings	4
2.2	Pre-Data and Non-Data Approaches to CLV Estimation	5
2.3	Traditional Approaches to CLV Estimation	6
2.3.1	RFM Scoring and Heuristic Segmentation	6
2.3.2	Deterministic CLV Formulas	7
2.4	Machine-Learning Approaches to CLV	9
2.5	Probabilistic Models for Customer-Base Analysis	10
2.6	Foundations for the Hierarchical and Decision-Theoretic Extensions	11
2.6.1	Hierarchical Models and Borrowing Strength	11
2.6.2	Bayesian Decision Theory and Uncertainty-Aware Targeting	12
2.7	Research Gap and Original Contribution	12
2.8	Summary of Theoretical Findings	13

3 Bayesian Probabilistic CLV Models 15

3.1	Bayesian Inference Fundamentals	15
3.1.1	Markov Chain Monte Carlo and the NUTS Sampler	16
3.2	Probability Distributions for Modelling Customer Behaviour	17
3.2.1	Purchase Timing and Counts: Poisson, Exponential, and Gamma	17
3.2.2	Dropout: Geometric and Beta	18
3.2.3	Monetary Value: the Gamma-Gamma Compound	19
3.2.4	The Building Blocks in the Study Data	19
3.2.5	Prior Distributions	21
3.3	The BG/NBD Model for Purchase Frequency	22
3.3.1	Generative Story	22
3.3.2	Key Model Outputs	22
3.3.3	Model Assumptions and Limitations	24
3.4	The Gamma-Gamma Model for Monetary Value	25
3.4.1	Model Specification	25
3.4.2	Key Assumption: Independence of Frequency and Monetary Value	25
3.4.3	CLV Computation	25
3.5	Hierarchical Extensions: Partial Pooling Across Segments	26
3.5.1	Partial Pooling vs. Alternatives	27
3.5.2	Implementation Considerations	27
3.6	Decision Theory Under Uncertainty	28
3.6.1	The Targeting Problem	28
3.6.2	Expected Improvement Over Point Estimates	29
3.7	Summary of Theoretical Findings and Hypotheses	30
4	Methodology	31
4.1	Research Design	31
4.2	Dataset	31
4.3	Data Preparation	32
4.3.1	Cleaning	32
4.3.2	Temporal Split	32

4.3.3	Customer-Level Aggregation	32
4.3.4	Country Segmentation	34
4.4	Exploratory Findings	34
4.5	Model Specification	35
4.6	Bayesian Inference	36
4.7	Baseline Models	36
4.8	Prediction	37
4.9	Evaluation Protocol	37
4.10	Implementation and Reproducibility	39
5	Results	40
5.1	MCMC Convergence and Diagnostics	40
5.2	H1: Predictive Accuracy and Calibrated Uncertainty	42
5.2.1	Point and Ranking Accuracy	42
5.2.2	Uncertainty Calibration	44
5.2.3	Activity Prediction	45
5.2.4	A Worked Example: One Customer End to End	45
5.3	H2: Hierarchical Partial Pooling Across Segments	46
5.4	H3: Decision-Theoretic Targeting	49
5.5	Summary of Hypothesis Tests	52
6	Discussion and Conclusion	53
6.1	Summary of Findings	53
6.2	Theoretical Contributions	54
6.3	Managerial Implications	54
6.4	Limitations	55
6.4.1	Modelling Assumptions	55
6.4.2	Evaluation Design	55
6.4.3	Scope of the CLV Construct	56
6.5	Generalisability	56

6.6	Future Work	56
6.7	Concluding Remarks	57
	References	58

List of Figures

3.1	Graphical model of the BG/NBD + Gamma-Gamma framework	23
3.2	Buy-till-you-die timeline	23
3.3	Empirical recency versus observation period (T) for the calibration sample. Every customer necessarily lies on or below the diagonal recency = T ; the spread of points away from the diagonal encodes the latent-activity signal the BG/NBD model exploits.	24
3.4	Correlation structure among the RFM summary statistics in the calibration data. The frequency–monetary association is the quantity that must be near zero for the Gamma-Gamma independence assumption to hold.	26
3.5	Hierarchical partial pooling	27
3.6	Customer counts by country segment in the calibration data. The pronounced imbalance, a dominant UK segment alongside several small markets, is precisely the regime in which partial pooling is expected to help.	28
3.7	Decision theory under uncertainty	29
4.1	Data-cleaning waterfall: the cumulative effect of each filtering step on the row count, from 1,067,371 raw line items to 802,679 retained transactions. .	32
4.2	Monthly transaction revenue across the full observation window. The calibration/holdout boundary is fixed at 1 March 2011.	33
4.3	Distribution of repeat-purchase frequency in the calibration sample. The mass at low frequencies and the long right tail are characteristic of non-contractual retail.	34
4.4	Distribution of average transaction value among repeat purchasers. The right skew motivates the Gamma-Gamma specification.	35
5.1	Split- $\hat{R}$ convergence diagnostics for the pooled BG/NBD model. The dashed line marks the 1.01 threshold.	40
5.2	Posterior trace plots for the pooled BG/NBD population parameters (r, α, a, b). Well-mixed, stationary chains indicate reliable sampling.	41
5.3	Pairwise posterior marginals for the BG/NBD parameters, with divergences highlighted. Strong correlations would motivate reparameterisation.	41

5.4	Split- $\hat{R}$ for the hierarchical BG/NBD model under the non-centred parameterisation.	43
5.5	Decile calibration: mean predicted versus mean actual holdout transactions for each model. The diagonal denotes perfect calibration.	44
5.6	Posterior predictive distributions of holdout transactions for a sample of customers (violins), with realised values overlaid. The spread of each violin is the model's per-customer uncertainty.	44
5.7	Distribution of $P(\text{alive})$ overall and by frequency group. Higher-frequency customers concentrate at higher $P(\text{alive})$	45
5.8	Two posteriors for one customer (ID 14817). The narrow posterior of $E[\text{CLV}]$ (orange) reflects parameter uncertainty only and lies entirely above the £600 cost, giving $P(\text{CLV} > \text{£}600) = 1.00$. The posterior predictive of realised CLV (blue) is far wider, placing $\sim 23\%$ of its mass below the cost, so the same probability computed correctly is 0.77. The realised holdout value (£1010) falls well inside the predictive distribution.	47
5.9	Per-country transaction MAE by model. Differences concentrated in the small segments (Germany, France, Other) test the partial-pooling hypothesis.	48
5.10	Per-segment posterior of the BG/NBD parameter r (90% HDI), with the global mean marked. Small segments are shrunk toward the global estimate.	48
5.11	Decision-theoretic targeting for the Bayesian model at the headline cost of £600. Left: cumulative net value by strategy and depth. Right: the net-value difference between the posterior-probability rule and the point estimate.	50
5.12	Distribution of the per-customer targeting probability $P(\text{CLV} > \text{£}600)$ across all 4,522 customers. (a) Computed from the posterior of $E[\text{CLV}]$, the probability saturates: 96% of customers sit at 0 or 1. (b) Computed from the posterior predictive of realised CLV, it is graded and usable as a decision score. This is the saturation artefact that makes the expectation-based rule fail for mechanical reasons.	50
5.13	Cumulative-gain and lift curves for all models. Steeper early rise indicates better identification of high-value customers.	51

List of Tables

2.1	The evolution of CLV estimation, from deterministic heuristics to generative Bayesian models. Only the probabilistic-generative tradition produces calibrated, distributional value estimates from raw transaction data.	8
3.1	Probability distributions used in the generative CLV models and the behavioural role of each.	19
3.2	Estimation strategies for multi-segment customer data.	27
3.3	Testable hypotheses evaluated in the empirical analysis.	30
4.1	Mapping of evaluation metrics to research hypotheses.	37
4.2	Evaluation metrics at a glance: plain-language intuition and why each is used.	38
5.1	Model comparison across transaction and CLV prediction metrics.	42
5.2	Posterior predictive credible interval coverage and CRPS.	44
5.3	P(alive) evaluation: binary activity prediction metrics.	45
5.4	Per-country transaction-prediction MAE by model (RQ2/H2).	47
5.5	Decision-theoretic targeting (BG/NBD (Bayesian)): posterior-probability vs point-estimate vs oracle net value, swept over intervention cost.	49
5.6	Risk-sensitive targeting metrics at the headline cost of £600 and a targeting depth of 20%. <i>Hit rate</i> is the fraction of targeted customers whose realised holdout value exceeds the intervention cost; <i>wasted spend</i> is the cumulative shortfall on targeted customers who fall below it. The point-estimate rule and the posterior- <i>predictive</i> probability rule are close on both; the probability rule computed from the posterior of <i>expected</i> CLV instead collapses, illustrating the saturation artefact.	51
5.7	Summary of hypothesis-test outcomes.	52

List of Abbreviations

Abbrev.	Expansion
AUC	Area under the ROC curve
BFMI	Bayesian fraction of missing information
BG/NBD	Beta-Geometric/Negative-Binomial-Distribution model
BTYD	Buy-till-you-die (model family)
CLV	Customer lifetime value
CRPS	Continuous ranked probability score
ESS	Effective sample size
HDI	Highest-density interval
HMC	Hamiltonian Monte Carlo
LOO-CV	Leave-one-out cross-validation
MAE	Mean absolute error
MAPE	Mean absolute percentage error
MCMC	Markov chain Monte Carlo
MLE	Maximum-likelihood estimation
NBD	Negative binomial distribution
NDCG	Normalised discounted cumulative gain
NUTS	No-U-Turn Sampler
RFM	Recency, frequency, monetary (value)
RMSE	Root mean squared error
RQ / H	Research question / hypothesis
WAIC	Widely applicable information criterion

Notation

Symbol	Meaning
λ	Individual customer's Poisson purchase rate (per week)
p	Individual dropout probability after a purchase (BG/NBD)
r, α	Shape and rate of the Gamma distribution mixing λ across customers
a, b	Shape parameters of the Beta distribution mixing p across customers
x	Frequency: number of <i>repeat</i> purchases in the calibration window
t_x	Recency: time from first to last calibration purchase (weeks)
T	Length of the customer's observation window (weeks)
m	Mean monetary value per repeat transaction
p, q, γ	Population parameters of the Gamma-Gamma spend model
v	Individual customer's Gamma-Gamma spend-rate parameter
$P(\text{alive})$	Posterior probability a customer is still active at T
c	Per-customer intervention (targeting) cost
τ	Posterior-probability threshold in the targeting rule
d	Discount rate in the deterministic CLV formula
θ	Generic parameter vector
$\mathcal{D}$	Observed data
$\hat{R}$	Gelman–Rubin convergence diagnostic

Chapter 1

Introduction

1.1 Background and Motivation

In the modern digital economy, the ability to anticipate future customer behaviour is among the most commercially valuable capabilities a business can possess. Whether allocating limited marketing budgets, designing retention programmes, or pricing acquisition channels, managers routinely face decisions whose quality depends on how well they can estimate the long-term value of individual customers. Customer Lifetime Value (CLV), the expected net present value of all future cash flows from a customer relationship, provides a theoretically grounded and practically actionable foundation for these decisions (Gupta et al., 2006). As Blattberg and Deighton (1996) argued, customers are financial assets whose value can and should be measured with the same rigour applied to capital investments; CLV formalises this intuition into a computable quantity.

Despite its conceptual clarity, accurately estimating CLV in practice is considerably more complex than the underlying formula implies. In non-contractual contexts, which are prevalent in e-commerce, where customers transact sporadically without formal subscriptions, organisations cannot directly ascertain when a customer has permanently churned. At any given time, a customer who has not transacted for several months may simply be in an extended inter-purchase interval or may have silently defected. This latent attrition issue implies that naive extrapolation of historical purchase frequencies will systematically overestimate the value of customers who have, in reality, discontinued their relationship. Consequently, the central challenge extends beyond predictive accuracy to encompass probabilistic reasoning under uncertainty about an inherently unobservable state (Fader & Hardie, 2009).

1.2 Understanding Customer Lifetime Value: Concept and Historical Development

CLV is a financial construct: the anticipated net present value of all future cash flows from a customer relationship. Rather than valuing isolated transactions, it prompts firms to estimate each customer's future spending potential, the likely duration of the relationship, and the uncertainty attached to both, and it thereby informs the allocation of marketing budgets, the design of retention initiatives, and the assessment of customer-base health. Blattberg and Deighton (1996) argued that customers should be valued as financial assets with the same

rigour applied to capital investments, and Gupta et al. (2006) showed that even modest gains in retention yield disproportionately large increases in firm value, a sensitivity that makes incremental improvements in CLV-model accuracy commercially significant.

The concept's development reflects marketing's shift from a transaction-based orientation, judged by aggregate indicators such as market share, to a relationship-centric one in which the enduring profitability of individual customers is central. Two empirical regularities drove that shift: customer value is heavily concentrated in a small subset of customers (Storbacka, 1997), and small reductions in defection produce large profitability gains (Reichheld & Sasser, 1990). The methodological lineage that followed, from the practitioner RFM framework (Hughes, 1994) through the net-present-value formalisation of Dwyer (1989) to the probabilistic "buy-till-you-die" models of Schmittlein et al. (1987) and Fader et al. (2005) and, more recently, machine-learning predictors (Chamberlain et al., 2017), is reviewed systematically in Chapter 2. This thesis positions itself within the probabilistic tradition, contending that Bayesian inference applied to generative CLV models offers a particularly rigorous framework for estimating individual customer value under uncertainty, and validates this claim empirically.

1.3 Customer Lifetime Value in E-Commerce and Modern Business

In e-commerce the economics of digital retail make customer-level CLV measurement essential rather than merely advantageous. Online retailers face explicit, quantifiable per-customer acquisition costs (paid search, affiliate fees, social-media campaigns) that must be justified by anticipated long-term returns, and in many sectors initial purchases do not cover them. Overestimating CLV leads to excessive bidding for new customers and value destruction; underestimating it cedes market share to more aggressive competitors (Gupta et al., 2006). Retention marketing poses the same problem in a different guise: without a forward-looking value estimate, firms cannot judge whether a retention initiative's cost is justified, or distinguish customers meriting resource-intensive intervention from those for whom low-cost outreach suffices. Reinartz and Kumar (2000) showed empirically that the link between tenure and profitability is weaker than commonly assumed, with many long-term customers unprofitable and some short-term customers contributing outsized value, which argues for individual-level CLV models over broad retention benchmarks.

CLV has also become central to firm-level analysis: McCarthy and Fader (2018) demonstrated that valuing a company by aggregating individual customer lifetime values often yields more precise and timely valuations than conventional accounting methods, allowing stakeholders to distinguish unsustainable growth from rational investment in customer assets. The UCI Online Retail II dataset used in this thesis exemplifies the non-contractual, episodic purchasing setting in which robust CLV estimation is simultaneously most challenging and

most consequential.

1.4 Research Questions and Hypotheses

This thesis is organised around three research questions (RQ) and their associated hypotheses (H). Each question addresses one of the three gaps identified in the literature review (Section 2.7), is motivated theoretically in Chapter 3, is operationalised by a dedicated element of the empirical design (Chapter 4), and is tested in Chapter 5:

RQ1 / H1: Accuracy and calibrated uncertainty.

Does the Bayesian BG/NBD + Gamma-Gamma model achieve competitive or superior out-of-sample accuracy relative to RFM-heuristic and XGBoost baselines, while additionally delivering well-calibrated uncertainty estimates?

RQ2 / H2: Hierarchical partial pooling.

Does a hierarchical extension that partially pools information across country segments reduce prediction error for small-market customers relative to complete-pooling and no-pooling alternatives?

RQ3 / H3: Decision-theoretic targeting.

Does targeting customers by the posterior probability $P(\text{CLV} > c)$ produce higher cumulative realised value than targeting by point-estimate CLV rankings?

1.5 Structure of the Thesis

Chapter 2 reviews the theoretical foundations of CLV estimation, tracing the progression from pre-data heuristics through deterministic formulas and RFM scoring to modern machine-learning approaches, and isolating the gaps that motivate a probabilistic treatment. Chapter 3 develops the Bayesian probabilistic framework that constitutes the methodological core of the thesis: Bayesian inference and MCMC, the BG/NBD and Gamma-Gamma generative models, their hierarchical extension via partial pooling, and the decision-theoretic targeting framework. The chapter closes by formalising the three hypotheses above.

Chapter 2

Theoretical Foundations of Customer Lifetime Value

2.1 Customer Lifetime Value: Concept

CLV represents the total net present value of all future cash flows attributed to a customer over the entirety of their relationship with a firm. As a forward-looking metric, CLV shifts managerial attention from short-term transactional revenue toward long-term customer equity, enabling firms to make economically rational decisions about acquisition spending, retention investment, and resource allocation across their customer base (Gupta et al., 2006).

The concept of CLV emerged from the broader customer-equity literature pioneered by Blattberg and Deighton (1996), who argued that customers should be treated as financial assets whose value can be measured, managed, and optimised. Rust et al. (2000) formalised this perspective into the customer-equity framework, decomposing total firm value into the sum of individual customer lifetime values. This conceptual shift has profound implications for marketing practice: rather than evaluating campaigns solely on immediate return, marketers can assess investments based on their impact on the long-term value of the customer base. Formally, CLV can be expressed as

$$\text{CLV} = \sum_{t=1}^T \frac{m_t r_t}{(1+d)^t}, \quad (2.1)$$

where m_t is the expected margin in period t , r_t is the retention probability, d is the discount rate, and T is the time horizon. This deterministic formulation, while intuitive, requires point estimates of future margins and retention rates, introducing substantial uncertainty that the formula itself does not capture. This limitation motivates the probabilistic approaches that form the core of this thesis.

2.1.1 Contractual vs. Non-Contractual Settings

A critical distinction in the CLV literature is between contractual and non-contractual business settings (Fader & Hardie, 2009). In contractual settings (subscription services, insurance policies, or mobile-phone contracts), the firm observes when a customer becomes inactive, because the customer explicitly cancels or fails to renew. The modelling task reduces to predicting *when* customers will terminate, and standard survival-analysis methods apply

directly.

Non-contractual settings, typical of most retail and e-commerce environments, present a fundamentally harder problem. Customers do not announce their departure; they simply stop purchasing. At any point, the firm cannot distinguish with certainty between a customer who has permanently defected and one who is merely experiencing a long inter-purchase interval. This latent-attribution problem requires probabilistic models that jointly estimate the customer's transaction rate and their unobserved dropout process, making it a natural domain for Bayesian inference. This thesis focuses exclusively on the non-contractual setting, as it represents the more challenging and practically common scenario in e-commerce. The UCI Online Retail II dataset used in the empirical part reflects this setting: customers purchase episodically with no formal contract, and the firm must infer activity status from observed purchase patterns.

2.2 Pre-Data and Non-Data Approaches to CLV Estimation

Prior to the advent of sophisticated transaction-level data and the computational power required to analyse it, organisations relied on a combination of judgment-based and formulaic strategies to estimate CLV. Although these early approaches lacked the precision of contemporary models, they laid the groundwork for the conceptual frameworks that underpin current CLV methodologies. The most prevalent legacy technique was the deterministic CLV formula, which calculated expected customer value by multiplying average order value, purchase frequency, and an estimated customer lifespan, then discounting to present value using a managerially determined rate. This method, detailed in the direct-marketing literature by Hughes (1994) and Stone (1995), was valued for its clarity: managers could trace the logic from inputs to outcome, which made it effective for justifying budgets or informing strategic discussions. Its core drawback lay in its reliance on assumed or historically averaged inputs, offering no means to capture genuine differences among individual customers, precisely the nuance needed for effective targeting. Applying the same expected lifespan to both a frequent purchaser and a one-time buyer, with only average order value as a differentiator, conflated distinct customer types and led to aggregate estimates that might be reasonable in total but largely uninformative at the individual level.

A related method, cohort-based value estimation, grouped customers acquired in the same timeframe and tracked their combined revenue over subsequent periods, producing a survival curve that firms could extrapolate to predict the remaining value of active members. Dwyer (1989) identified this as one of the earliest formal CLV treatments. While it leveraged actual retention data, improving on the single-formula method, the approach still operated at an aggregate level and required untestable assumptions about the projected survival curve beyond observed data.

An incremental move toward individualisation came with RFM scoring, which derived composite scores from customers' transaction histories, allowing companies to identify high-value customers for targeted marketing without building formal statistical models. Popularised by Hughes (1994) and Cullinan (1977) in direct mail, RFM remains widespread in practice. Yet its limitations are clear: it reflects only past behaviour, lacks any measure of confidence or predictive reliability for individual rankings, and cannot distinguish genuinely promising customers from those who simply made an unusually large recent purchase.

Expert judgment and managerial intuition represented another tier of pre-data CLV estimation, especially in business-to-business contexts where transaction volumes were low and individual insights could be obtained directly by account teams. Sales teams' estimates of renewal probabilities and expected contract values served as informal proxies for CLV, which management aggregated into revenue forecasts. While subjective, these estimates often incorporated relationship-specific intelligence inaccessible to algorithmic models. However, Lovallo and Kahneman (2003) and others document systematic biases (overconfidence in renewals, recency effects, and the inflation of estimates to protect internal resources) that ultimately motivated the transition to data-driven CLV models as more granular data became available.

The unifying weakness of these pre-data approaches is that each delivered a single point estimate per customer or cohort, with no associated measure of confidence, obscuring the difference between reliably high-value customers and those whose estimates masked substantial uncertainty. Introducing calibrated uncertainty is a central purpose of the Bayesian models developed in this thesis.

2.3 Traditional Approaches to CLV Estimation

2.3.1 RFM Scoring and Heuristic Segmentation

The Recency, Frequency, and Monetary (RFM) framework remains one of the most prevalent methods for assessing customer value in contemporary marketing (Hughes, 2011). It evaluates customers along three behavioural dimensions: the time elapsed since their last purchase (Recency), the number of transactions (Frequency), and cumulative expenditure (Monetary). Each customer is ranked within the population on these dimensions, usually through quintile scoring, and the combined RFM score is used as an indicative measure of individual customer value.

RFM's enduring popularity can be attributed to its straightforward implementation, clarity, and modest data prerequisites: it relies solely on transactional records containing customer IDs, purchase dates, and purchase amounts. Its computational ease makes it accessible to

practitioners without technical backgrounds, and its utility has been demonstrated repeatedly over decades of use in direct marketing (Blattberg et al., 2008). Nevertheless, several core limitations become apparent when RFM is employed to estimate CLV:

- **Lack of predictive capability.** While RFM scores efficiently summarise historical behaviour, they offer no explicit means to forecast future purchases or anticipated spending. A customer currently exhibiting a high RFM score may soon become inactive, yet this framework cannot identify such impending attrition.
- **Absence of uncertainty estimation.** RFM generates a single, deterministic value per customer without any indication of statistical confidence. It is therefore impossible to discern whether a high score reliably reflects sustained engagement or is simply due to a recent, anomalous transaction.
- **Artificial discretisation.** Dividing customers into quintiles imposes rigid boundaries that may not reflect meaningful differences. Two individuals with almost identical recency can be assigned disparate scores, while substantially different behaviours may be masked within a single category.
- **Lack of an underlying generative model.** RFM offers no explicit account of the process driving customer behaviour. In its absence, it is difficult to extrapolate future actions, systematically integrate prior information, or compare alternative approaches on a unified statistical basis.

Owing to these shortcomings, there is a clear impetus for adopting model-based CLV techniques that deliver probabilistic forecasts alongside robust measures of uncertainty.

2.3.2 Deterministic CLV Formulas

Beyond RFM heuristics, several deterministic formulas have been proposed for CLV estimation. The most common deterministic approach projects customer value using average purchase frequency and average order value over an assumed customer lifespan, discounted to present value. Variants include the simple CLV formula popularised by Gupta and Lehmann (2003), which assumes constant margin and retention rates, and the migration model of Dwyer (1997), which tracks customers through discrete value states. While these formulas provide closed-form solutions that are easy to implement and communicate, they share a common weakness: they require aggregate-level parameter assumptions (e.g., a single retention rate applied uniformly) that obscure the substantial heterogeneity present in real customer populations. A firm with an average annual retention rate of 80% may contain segments with 95% retention and others with 50%, and a uniform-rate formula would produce misleading CLV estimates for both groups.

Table 2.1: The evolution of CLV estimation, from deterministic heuristics to generative Bayesian models. Only the probabilistic-generative tradition produces calibrated, distributional value estimates from raw transaction data.

Approach	Core idea	Individual-level	Uncertainty
Deterministic formula	Average order value $\times$ frequency $\times$ lifespan, discounted	No	No
Cohort survival	Track acquired cohorts; extrapolate a survival curve	No	No
RFM scoring	Rank customers by recency, frequency, monetary quintiles	Partial	No
Probabilistic (buy-till-you-die)	Generative model of purchasing and latent dropout	Yes	Point (MLE)
Machine learning	Discriminative regression on engineered features	Yes	Post-hoc
Bayesian generative	Full posterior over BG/NBD + Gamma-Gamma parameters	Yes	Calibrated

Note: author's own compilation, synthesising Blattberg & Deighton (1996), Schmittlein et al. (1987), Fader et al. (2005), Chamberlain et al. (2017), and related literature.

2.4 Machine-Learning Approaches to CLV

The growing availability of granular customer data has motivated the application of supervised machine-learning methods to CLV prediction. Typical implementations frame CLV as a regression problem: features are engineered from historical transaction data (often RFM-based, augmented with additional behavioural or demographic variables), and models such as gradient-boosted trees (XGBoost), random forests, or neural networks are trained to predict future spend or purchase counts over a fixed horizon (Chamberlain et al., 2017). Machine-learning approaches offer several advantages over traditional methods:

- **Flexible function approximation.** Tree-based models can capture complex, non-linear relationships between features and CLV without requiring explicit functional-form assumptions.
- **Feature richness.** ML models can incorporate a wide range of features beyond RFM, including browsing behaviour, product-category preferences, seasonality indicators, and customer demographics.
- **Strong point-prediction accuracy.** On well-specified tasks with sufficient training data, ML models frequently achieve lower mean absolute error or root mean squared error than parametric models.

However, ML approaches to CLV carry significant limitations that constrain their usefulness for decision-making:

- **No generative model.** Like RFM, ML regression models are discriminative rather than generative: they learn a mapping from features to targets without specifying the underlying data-generating process. This makes it difficult to reason about *why* predictions differ across customers or to simulate counterfactual scenarios (e.g., “what would this customer’s CLV be if we increased contact frequency?”).
- **Uncertainty is not native to standard pipelines.** Standard implementations of machine-learning models focus on point prediction and typically require additional techniques to obtain calibrated uncertainty. Dedicated methods do exist, including quantile regression, conformal prediction (Angelopoulos & Bates, 2023), NGBoost, deep ensembles, and Monte Carlo dropout, but in conventional pipelines these are layered on top of a point predictor rather than being intrinsic to the framework, and the resulting intervals are often poorly calibrated, particularly in the tails of the distribution where decisions matter most.
- **Dependence on feature engineering.** The quality of ML-based predictions depends heavily on manual feature engineering. Choosing the right aggregation windows, lag features, and transformations requires domain expertise that may not transfer across

business contexts. Probabilistic models, by contrast, operate directly on raw transaction data and derive behavioural parameters endogenously.

- **Fixed prediction horizon.** A model trained to predict “spend in the next 6 months” cannot be trivially reused for a 12- or 24-month horizon without retraining. Generative CLV models produce the full future transaction and spend distribution, from which any horizon can be derived.

These limitations do not make ML approaches unsuitable for CLV estimation, but they motivate the investigation of generative Bayesian models that address these gaps natively. In the empirical part of this thesis, an XGBoost baseline serves as a benchmark to contextualise the Bayesian model’s performance.

2.5 Probabilistic Models for Customer-Base Analysis

The limitations of heuristic and discriminative methods motivated a distinct modelling tradition built on explicit *generative* stories for customer behaviour. These “buy-till-you-die” (BTYD) models, originating with Schmittlein et al. (1987), treat each customer’s observed purchasing as the outcome of two unobserved processes: a transaction process governing how often an active customer buys, and a dropout process governing when the customer becomes permanently inactive. Because the dropout time is never observed in a non-contractual setting, the model infers the probability that each customer is still active from recency and frequency alone, directly confronting the latent-attrition problem that defeats deterministic formulas.

The canonical BTYD specification is the Pareto/NBD model of Schmittlein et al. (1987), which pairs a Poisson transaction process with an exponentially distributed lifetime and lets both rates vary across customers through gamma mixing distributions. The Pareto/NBD enjoys strong empirical support but is computationally awkward: its likelihood involves Gauss hypergeometric functions and numerous numerical evaluations, which historically hindered adoption. Fader et al. (2005) introduced the BG/NBD model as a more tractable alternative that replaces the continuous-time dropout assumption with a simpler rule, namely that a customer may become inactive only immediately after a purchase, governed by a beta-distributed dropout probability. BG/NBD reproduces the predictive performance of the Pareto/NBD at a fraction of the computational cost and has since become the de facto standard for non-contractual customer-base analysis (Fader & Hardie, 2009).

BTYD models in this tradition describe only the timing and number of transactions; they are silent about how much each transaction is worth. To obtain a monetary CLV, a transaction-count model is therefore combined with a separate spend model. The standard choice is the Gamma-Gamma model (Fader & Hardie, 2013), which represents each customer’s aver-

age transaction value as a draw from a gamma distribution whose own rate parameter varies across customers, and which assumes that average spend is independent of purchase frequency. This independence assumption is convenient because it allows the frequency and monetary components to be estimated separately and then multiplied, but it is an empirical assumption that must be verified on the data at hand. The BG/NBD + Gamma-Gamma pairing, and the conditions under which it is appropriate for the UCI Online Retail II data, are developed in detail in Chapter 3.

Recent research extends this tradition in two directions. On the model side, Bachmann et al. (2021) incorporate time-varying contextual factors (seasonality, holidays, marketing shocks) into latent-attribution models, relaxing the stationarity assumption that limits the classical specifications, and provide the CLVTools implementation now standard in applied work. On the methodological side, neural sequence models replace the parametric generative story altogether: Valendin et al. (2022) show that recurrent networks trained on raw transaction sequences can match or exceed Pareto/NBD-type models on aggregate forecasting tasks. Both strands, however, pursue predictive accuracy through richer model classes while leaving open the questions this thesis addresses: whether parameter and prediction uncertainty is *calibrated*, and whether the resulting distributions improve marketing *decisions*. These inferential and decision-theoretic questions apply equally to the richer models and are examined here on the tractable BG/NBD + Gamma-Gamma foundation, where they can be isolated cleanly.

A common feature of these probabilistic models, as conventionally applied, is that their population parameters are fitted by maximum likelihood, yielding point estimates with, at best, asymptotic standard errors. They also assume a single homogeneous population, fitting one set of parameters to all customers regardless of segment. The remainder of this thesis argues that both features can be improved: a fully Bayesian treatment propagates parameter uncertainty into every prediction, and a hierarchical extension relaxes the single-population assumption by partially pooling information across customer segments.

2.6 Foundations for the Hierarchical and Decision-Theoretic Extensions

2.6.1 Hierarchical Models and Borrowing Strength

When a customer base divides into segments of very unequal size, estimating behavioural parameters separately for each segment is unreliable for the small ones, whose few observations yield noisy estimates, whereas pooling all segments into a single model ignores genuine differences between them. Hierarchical (multilevel) models resolve this tension by treating segment-level parameters as draws from a common population distribution, so that each seg-

ment's estimate is shrunk toward the global mean by an amount that depends on how much data the segment provides (Gelman & Hill, 2006). The classical justification for this “borrowing of strength” is the result of Efron and Morris (1975), who showed that shrinkage estimators can dominate independent per-group estimates in mean-squared error when several related quantities are estimated jointly. In marketing specifically, Rossi et al. (2005) demonstrate that hierarchical Bayesian models are well suited to customer-level data, where individuals or segments differ yet share enough structure to inform one another. This logic motivates the hierarchical extension evaluated in this thesis (RQ2): country segments such as Germany and France contribute few customers relative to the dominant UK segment, and partial pooling is expected to stabilise their estimates without discarding the heterogeneity that distinguishes them.

2.6.2 Bayesian Decision Theory and Uncertainty-Aware Targeting

Predicting customer value is ultimately a means to an end: deciding where to direct finite marketing resources. Bayesian decision theory provides a formal framework for such choices under uncertainty, prescribing the action that maximises expected utility, or equivalently minimises expected loss, with respect to the full posterior distribution over the unknown quantities (Berger, 1985). This perspective bears directly on customer targeting, conventionally framed as selecting customers whose predicted value exceeds the cost of an intervention (Gupta et al., 2006; Venkatesan & Kumar, 2004). When targeting is based on point predictions alone, two customers with the same predicted value are treated identically even if one estimate is precise and the other highly uncertain. A decision rule defined on the posterior, by contrast, can account for that difference, for instance by targeting a customer only when the probability that their value exceeds the intervention cost is sufficiently high. Whether this uncertainty-aware rule yields better realised outcomes than point-estimate targeting is the question posed by RQ3.

2.7 Research Gap and Original Contribution

Three specific gaps emerge from the review above. Each maps onto one research question and onto a dedicated element of the empirical design, so that the chain from gap to hypothesis to method is explicit.

Gap 1 (inference): uncertainty is produced but rarely shown to be calibrated. The BTYD tradition is generative but conventionally estimated by maximum likelihood, providing at best asymptotic parameter uncertainty (Section 2.5); machine-learning approaches require uncertainty to be retrofitted, with calibration not guaranteed (Section 2.4; Angelopoulos and Bates, 2023). What is missing is empirical evidence that a fully Bayesian treatment

delivers competitive accuracy *and* demonstrably calibrated uncertainty on real transaction data. This is RQ1/H1, operationalised through the coverage, CRPS, and accuracy protocol of Section 4.9.

Gap 2 (structure): segment imbalance is pervasive but the pooling question is rarely tested. Hierarchical Bayesian methods are well established in marketing (Rossi et al., 2005), yet applied BTYD studies almost always fit a single homogeneous population, and the three-way comparison of complete pooling, partial pooling, and no pooling is seldom carried out on customer-base data with realistic segment imbalance. This is RQ2/H2, operationalised as the three-way pooling comparison of Sections 4.5 and 5.3.

Gap 3 (decisions): posteriors are computed but not used. Where posterior distributions are available, they are seldom carried through to the targeting decision itself; the decision-theoretic framework (Berger, 1985; Venkatesan & Kumar, 2004) predicts that they should matter precisely when estimates are uncertain near the decision boundary. This is RQ3/H3, operationalised as the targeting simulation of Section 4.9.

The original contribution of the thesis is correspondingly threefold: (i) an end-to-end Bayesian implementation of the BG/NBD + Gamma-Gamma framework estimated by Hamiltonian Monte Carlo, evaluated calibration-first rather than accuracy-only; (ii) a hierarchical extension tested against *both* pooling extremes rather than a single benchmark; and (iii) a decision-theoretic evaluation that, beyond testing H3, identifies a general and practically consequential pitfall: decision rules must be computed from the posterior *pre-dictive* distribution, not from the posterior of an expectation, which saturates and degrades the decision (Section 5.4). Recent neural and covariate-rich extensions of customer-base analysis (Bachmann et al., 2021; Valendin et al., 2022) pursue predictive accuracy through richer model classes; the inferential and decision-theoretic contributions made here are complementary and remain applicable to those models.

2.8 Summary of Theoretical Findings

Three findings from this chapter carry forward:

- Non-contractual CLV requires probabilistic modelling because attrition is latent; deterministic formulas and RFM heuristics cannot address it.
- Neither RFM nor standard machine-learning pipelines provides native, calibrated uncertainty, limiting their usefulness for risk-aware decisions.
- The “buy-till-you-die” tradition does model timing, dropout, and monetary value jointly, but is conventionally estimated by maximum likelihood, yielding limited information about parameter uncertainty. This motivates the Bayesian treatment

introduced next.

Chapter 3

Bayesian Probabilistic CLV Models

This chapter provides an overview of the Bayesian probabilistic approach to predicting CLV, the central methodological focus of this thesis. The discussion begins by outlining essential Bayesian principles, introduces the probability distributions from which the generative models are built and the priors placed on their parameters, then progresses to the specific generative models used in the empirical analysis, and concludes by examining the hierarchical extensions and decision-theoretic applications that constitute the thesis's main contributions.

3.1 Bayesian Inference Fundamentals

Bayesian inference provides a principled framework for updating beliefs about unknown parameters in light of observed data. Given a model with parameters θ , observed data $\mathcal{D}$, and a prior distribution $p(\theta)$ encoding beliefs before observing the data, Bayes' theorem yields the posterior distribution

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta) p(\theta)}{p(\mathcal{D})} \propto p(\mathcal{D} \mid \theta) p(\theta), \quad (3.1)$$

where $p(\mathcal{D} \mid \theta)$ is the likelihood and the normalising constant $p(\mathcal{D})$ (the marginal likelihood or evidence) ensures the posterior integrates to one. The posterior integrates the information provided by both the data and the prior, summarising all that is known about the parameters and underpinning subsequent inference and prediction (Gelman et al., 2013). Three properties of Bayesian inference are particularly relevant for CLV modelling:

- **Comprehensive uncertainty quantification.** Rather than a single value, the posterior yields a full probability distribution. Every subsequent estimate (future transactions, expected revenue, or CLV) reflects the inherent uncertainty, with credible intervals emerging naturally from the analysis.
- **Integration of prior knowledge.** Practitioners can embed domain expertise through informative priors (for example, known retention rates). This stabilises estimates, particularly for customers with limited data, while weakly informative priors ensure results are not overly influenced when substantial data is available.
- **Rigorous model comparison.** Bayesian approaches facilitate principled model selection through criteria such as the Widely Applicable Information Criterion (WAIC)

or Leave-One-Out Cross-Validation (LOO-CV), enabling robust comparison between alternative specifications (Vehtari et al., 2017).

3.1.1 Markov Chain Monte Carlo and the NUTS Sampler

For the majority of practically relevant models, including those in this thesis, it is not feasible to compute the posterior analytically. Contemporary Bayesian inference therefore typically employs Markov Chain Monte Carlo (MCMC) methods, which generate posterior samples that can approximate any desired summary, such as means, quantiles, or credible intervals. The empirical work here employs the No-U-Turn Sampler (NUTS), an adaptive form of Hamiltonian Monte Carlo (HMC) available within the PyMC probabilistic-programming library (Hoffman & Gelman, 2014; Salvatier et al., 2016).

HMC explores the posterior by exploiting its gradient rather than by blind random-walk proposals. The parameter vector is augmented with an auxiliary momentum variable, and the negative log-posterior is treated as a potential-energy surface: each iteration draws a fresh momentum, then simulates the resulting frictionless motion across the surface with a discrete *leapfrog* integrator whose steps follow the gradient of the log-posterior. Because the simulated trajectory approximately conserves total energy, it can travel far from the current point and still be accepted with high probability, producing distant, weakly correlated samples where random-walk Metropolis-Hastings would diffuse slowly. NUTS removes the need to hand-tune the trajectory length by extending each trajectory, through repeated doubling, until it begins to turn back on itself, and its warm-up phase adapts the integrator step size toward a target acceptance rate while estimating the posterior’s scale along each dimension (Hoffman & Gelman, 2014).

This mechanism explains both the efficiency of NUTS and its characteristic failure modes, several of which proved directly relevant to this thesis. The integrator’s step size is dictated by the sharpest curvature of the posterior anywhere along a trajectory; geometries with strong parameter correlations (ridges), funnels, or wildly different scales across dimensions therefore force tiny steps and very long trajectories, and sampling cost explodes even for models with only a handful of parameters. The *numerical formulation* of the likelihood matters for the same reason: two algebraically equivalent expressions can present very different energy surfaces, and intermediate quantities that overflow, cancel catastrophically, or become non-finite corrupt the gradient and energy computations on which the trajectory depends. The practical levers, employed in Chapter 4, follow directly: choose a numerically stable form of the likelihood, place priors whose scales match the data (Section 3.2.5), and reparameterise where the geometry itself is pathological, as in the hierarchical funnel discussed in Section 3.5.2. Sections 4.5 and 4.6 document how these considerations shaped the estimated models.

The quality of MCMC inference is assessed through standard diagnostics reported in the empirical analysis:

- **Effective Sample Size (ESS).** Measures the number of effectively independent draws from the posterior, accounting for autocorrelation between successive samples. Higher ESS indicates more informative posterior summaries.
- $\hat{R}$ (**R-hat**). Compares within-chain and between-chain variance across multiple independent chains. Values close to 1.0 indicate convergence; values above 1.01 suggest potential non-convergence (Vehtari et al., 2021).
- **Bayesian Fraction of Missing Information (BFMI).** Assesses whether the sampler adequately explores the energy distribution. Low values indicate problematic exploration, often caused by funnel geometries or strong posterior correlations.
- **Divergences.** Occur when the NUTS sampler encounters regions of high curvature, signalling possible model misspecification or the need for reparameterisation.

3.2 Probability Distributions for Modelling Customer Behaviour

The generative models of this chapter are assembled from a small set of standard probability distributions, each selected because its mathematical form mirrors a specific behavioural assumption. The models share a common two-layer logic: an *individual-level* process distribution describes what a single customer does (how often they buy, when they defect, how much they spend), and a *population-level* mixing distribution describes how the parameters of that process vary across customers. This section introduces the building blocks and their roles; Sections 3.3 and 3.4 then show how the BG/NBD and Gamma-Gamma models assemble them.

3.2.1 Purchase Timing and Counts: Poisson, Exponential, and Gamma

The elementary assumption about purchasing is that an active customer buys “at random” at a constant underlying rate λ : purchase occasions arrive according to a Poisson process. The number of transactions $X(t)$ in a window of length t then follows a Poisson distribution,

$$P(X(t) = k \mid \lambda) = \frac{(\lambda t)^k e^{-\lambda t}}{k!}, \quad k = 0, 1, 2, \dots, \quad (3.2)$$

and the times between successive purchases are exponentially distributed with the same rate. The exponential distribution is *memoryless*: the time already elapsed since the last purchase carries no information about when the next one is due. This property is what makes the Poisson assumption appropriate for non-contractual retail purchasing, where there is no billing

cycle or scheduled renewal imposing a rhythm on transactions, and it is precisely the property that forces latent attrition to be modelled separately, since a long silence is equally consistent with a slow-but-active customer and a departed one.

A single rate λ for all customers would, however, contradict a robust empirical regularity: observed purchase counts are strongly *over-dispersed*, with variance far exceeding the mean, because customer bases mix frequent and occasional buyers. The standard remedy is to let the rate itself vary across customers according to a Gamma distribution, $\lambda \sim \text{Gamma}(r, \alpha)$, with shape r and rate α . The Gamma distribution is the natural mixing choice for a positive rate: it is flexibly right-skewed, spans shapes from strongly peaked to near-exponential, and is conjugate to the Poisson likelihood, which keeps the mixture analytically tractable. Integrating the individual rate out of Equation (3.2) yields the negative binomial distribution (NBD) for the marginal purchase count,

$$\begin{aligned} P(X(t) = k) &= \int_0^\infty P(X(t) = k \mid \lambda) \text{Gamma}(\lambda \mid r, \alpha) d\lambda \\ &= \frac{\Gamma(r+k)}{\Gamma(r) k!} \left(\frac{\alpha}{\alpha+t} \right)^r \left(\frac{t}{\alpha+t} \right)^k, \end{aligned} \quad (3.3)$$

whose variance exceeds its mean and which therefore reproduces the over-dispersion the pure Poisson model cannot. This Poisson-Gamma pair has a long history as a purchase-incidence model in marketing and supplies the “NBD” component of both the Pareto/NBD and the BG/NBD (Fader et al., 2005; Schmittlein et al., 1987).

3.2.2 Dropout: Geometric and Beta

Latent attrition requires a second process. The BG/NBD models dropout as a Bernoulli trial after every purchase: with probability p the customer becomes permanently inactive, with probability $1 - p$ they continue. The number of purchases a customer makes before defecting is then geometrically distributed, so the probability of surviving j purchases, $(1 - p)^j$, declines geometrically. The geometric distribution is the discrete counterpart of the exponential lifetime assumed by the Pareto/NBD (Schmittlein et al., 1987); tying dropout opportunities to purchase occasions rather than to continuous time is exactly the simplification that makes the BG/NBD likelihood tractable (Fader et al., 2005).

As with purchase rates, dropout propensities differ across customers, and the mixing distribution must respect the parameter’s support: p is a probability confined to $[0, 1]$. The canonical choice is the Beta distribution, $p \sim \text{Beta}(a, b)$, with density proportional to $p^{a-1}(1 - p)^{b-1}$. Its two shape parameters accommodate loyal-dominated populations (mass near zero), churn-prone populations (mass near one), and polarised mixtures of both, and it is

conjugate to the Bernoulli dropout trials. The resulting Beta-Geometric mixture is the “BG” component of the model’s name.

3.2.3 Monetary Value: the Gamma-Gamma Compound

Transaction values in retail data are positive and heavily right-skewed: most baskets are modest while a minority are very large. The Gamma distribution is again the natural individual-level model, with each customer’s average transaction value drawn from a Gamma distribution whose scale is customer-specific. Placing a second Gamma distribution on that customer-specific parameter produces the Gamma-Gamma compound of Fader and Hardie (2013): the population-level marginal distribution of average spend is then heavier-tailed than any single Gamma, matching the pronounced skew of observed spending, while individual-level estimates are shrunk toward the population mean in proportion to how few transactions support them. Section 3.4 presents the full specification.

Table 3.1 summarises the building blocks and the behavioural assumption each encodes.

Table 3.1: Probability distributions used in the generative CLV models and the behavioural role of each.

Distribution	Level	Behavioural role
Poisson	Individual	Number of purchases by an active customer in a fixed window
Exponential	Individual	Memoryless inter-purchase times implied by the Poisson process
Geometric	Individual	Number of purchases before permanent dropout
Gamma	Population	Heterogeneity in purchase rates λ (yielding the NBD); heterogeneity in spend parameters (yielding the Gamma-Gamma)
Beta	Population	Heterogeneity in dropout probabilities p on $[0, 1]$

3.2.4 The Building Blocks in the Study Data

These distributional choices are not merely mathematically convenient; each corresponds to an empirical signature clearly visible in the UCI Online Retail II calibration sample analysed in this thesis (the dataset and its preparation are described in Chapter 4). Consider first the purchase counts. The 4,522 calibration customers average 3.78 repeat purchases over a mean observation window of 41.8 weeks, an average purchase rate of roughly 0.08 per week, or one purchase every 13 weeks. The variance of the repeat-purchase count, however, is 77.7, more than twenty times the mean, where a common-rate Poisson model would require variance equal to the mean. The frequency distribution (Figure 4.3) makes the source of

this over-dispersion plain: a large mass of customers with very few purchases coexists with a long right tail reaching 200 repeat transactions. This is exactly the regime the Gamma-mixed Poisson of Section 3.2.1 is built for, with the caveat that the most extreme frequencies belong to trade and wholesale buyers whose ordering is more regular than a Poisson process implies, a point taken up in the limitations of Chapter 6.

The dropout components are equally visible. Fully 1,493 customers, 33.0% of the calibration sample, made no repeat purchase at all. Under the generative story of Section 3.2.2, this mass at zero conflates two indistinguishable populations: customers who defected immediately after their first purchase and customers who remain active but buy slowly. Resolving that ambiguity probabilistically, rather than by an arbitrary recency cut-off, is precisely what the Beta-Geometric dropout structure combined with the memoryless transaction process enables, and it is why the $P(\text{alive})$ quantity of Section 3.3.2 exists at all. Fitting this sample requires the Beta heterogeneity distribution to accommodate both a churn-prone mass consistent with the one-third of one-time buyers and the loyal tail that generates the long purchase histories.

Monetary values show the skew the Gamma-Gamma compound presumes. Among the 3,029 repeat purchasers, the mean average transaction value is £385 against a median of £299, with a 95th percentile of £900 and a maximum above £8,400 (Figure 4.4): a distribution whose mean sits well above its median and whose tail spans an order of magnitude. The Gamma-Gamma independence assumption also looks tenable at first inspection, with a frequency–monetary correlation of only 0.13 among repeat purchasers; this condition is examined formally in Section 3.4.2. Finally, the pronounced segment imbalance of the sample, 4,152 UK customers against 74 in Germany and 54 in France, foreshadows why the hierarchical extension of Section 3.5, and with it the prior choices discussed next, matter for this dataset.

A customer to follow. To make the abstract machinery concrete, one customer is followed through the entire pipeline. Customer 14817 (United Kingdom) made five purchases in the calibration window, hence a frequency of $x = 4$ repeat transactions, spending an average of £526 per repeat order, with a recency of $t_x = 33.1$ weeks inside an observation window of $T = 49.2$ weeks. This is a moderately active buyer whose activity status is genuinely uncertain: the model must decide how much of their fifteen-week silence before the calibration cut-off signals slow-but-active purchasing rather than silent defection. Their raw profile (x, t_x, T, m) is the only input the BG/NBD + Gamma-Gamma model receives; Section 5.2.4 returns to them once the model is fitted, tracing how these four numbers become a $P(\text{alive})$ estimate, a full posterior distribution over their future value, and a concrete targeting decision.

3.2.5 Prior Distributions

A conceptual distinction matters here. The Gamma and Beta mixing distributions of the preceding subsections are *part of the likelihood*: they describe real heterogeneity across customers, and their parameters (r, α, a, b) are population quantities to be estimated. In the classical treatment these are fitted by maximum likelihood and the analysis stops there. The Bayesian treatment adds a further layer: *prior distributions* on the population parameters themselves, expressing uncertainty about the customer base as a whole before the data are observed. It is this layer that Bayes' theorem (Equation (3.1)) updates, and it is what allows parameter uncertainty to propagate into every downstream prediction.

This thesis employs *weakly informative* priors (Gelman et al., 2013): distributions that respect each parameter's support and rule out absurd values without materially constraining the plausible range. All population parameters of the BG/NBD and Gamma-Gamma models are positive, so they receive half-normal priors, whose support is $[0, \infty)$ and whose mass declines gently away from zero, providing mild regularisation while permitting large values if the data demand them. The critical practical question is the *scale* of these priors, because the parameters carry units: purchase-rate parameters are of order one on a weekly time scale, whereas the Gamma-Gamma scale parameter is denominated in pounds and must be of the order of hundreds. A prior that ignores these units can concentrate its mass far from the region the data support, distorting inference and degrading sampler performance. The appropriate diagnostic is the *prior predictive check* (Gelman et al., 2013): simulating data from the prior and likelihood alone and verifying that the implied behaviour (purchase counts, average spends) is of a plausible magnitude. Chapter 4 describes how the prior scales used here are set from summary statistics of the calibration data, a data-informed variant of this logic with ample precedent in marketing applications of hierarchical Bayesian methods (Rossi et al., 2005).

The hierarchical model of Section 3.5 adds one further prior choice. There, the four BG/NBD parameters vary by country segment, with each segment's parameters drawn (on the log scale, ensuring positivity) from a common population distribution. The *between-segment* standard deviation of that distribution receives its own half-normal hyperprior, and its scale governs how much pooling the model performs: a large scale approaches independent per-segment estimation, while a scale of zero collapses to complete pooling. The empirical implementation uses a deliberately tight scale, encoding the substantive expectation that purchasing behaviour differs only moderately across countries; the consequences of this choice for the H2 comparison, including its role as a caveat on the pooling results, are taken up in Chapters 4 and 6.

3.3 The BG/NBD Model for Purchase Frequency

The Beta-Geometric/Negative-Binomial-Distribution (BG/NBD) model, proposed by Fader et al. (2005), serves as a foundational probabilistic framework for modelling customer purchasing behaviour in non-contractual contexts. It simultaneously characterises two hidden processes: the rate at which customers make purchases while active, and the timing of their permanent disengagement.

3.3.1 Generative Story

The BG/NBD model assembles the distributions of Section 3.2 into the following data-generating process for each customer:

- **Transaction process.** While active, each customer makes purchases according to a Poisson process with individual-specific rate λ . The number of transactions in a period of length t follows a Poisson distribution with mean λt .
- **Dropout process.** Following each purchase, the customer faces an individualised probability p of becoming permanently inactive, akin to flipping a biased coin. This geometric dropout mechanism means the likelihood of continued activity decreases geometrically as additional transactions occur.
- **Heterogeneity in transaction rates.** Transaction rates vary across customers according to a Gamma distribution, $\lambda \sim \text{Gamma}(r, \alpha)$, where r is the shape and α the rate parameter.
- **Heterogeneity in dropout probabilities.** The dropout probability varies across individuals and is modelled as $p \sim \text{Beta}(a, b)$.

The four parameters (r, α, a, b) operate at the population level, shaping the distributions from which each customer's individual parameters are drawn. This hierarchical arrangement reflects the considerable behavioural diversity observed in real customer groups without introducing unnecessary complexity. Figure 3.1 depicts the full generative structure, including the Gamma-Gamma monetary component of Section 3.4, and Figure 3.2 illustrates the buy-till-you-die logic on a single customer's transaction timeline.

3.3.2 Key Model Outputs

Given a customer's observed history, summarised by frequency x (number of repeat transactions), recency t_x (time of last purchase), and observation period T , the BG/NBD model provides:




Figure 3.1: Plate representation of the joint BG/NBD + Gamma-Gamma model. Rounded rectangles are population-level hyperparameters; open circles are latent per-customer parameters (transaction rate λ_i , dropout probability p_i , spend rate v_i); shaded circles are observed data (repeat purchases x_i , recency $t_{x,i}$, observation window T_i , average value z_i). The plate denotes replication over all N customers.




Figure 3.2: The buy-till-you-die logic for one customer. Five transactions are observed; the model summarises them as frequency x , recency t_x , and observation window T . After the last purchase the customer's state is latent: the BG/NBD model infers $P(\text{alive})$ from how recently and how often they bought relative to T .

- **$P(\text{alive})$.** The posterior probability that a customer remains active at the end of the observation period, directly addressing the latent-attribution challenge inherent in non-contractual CLV estimation.
- **Conditional expected transactions.** $E[X(t) \mid x, t_x, T]$ denotes the expected number of purchases over a future period of length t , based on the customer's history. This output constitutes the transaction-count component of CLV.

Both outputs are computed as expectations over the posterior distributions of individual customer parameters, naturally integrating uncertainty about each customer's true transaction frequency and dropout probability. Figure 3.3 previews the empirical (recency, T) structure of the UCI Online Retail II calibration sample that these quantities are derived from.




Figure 3.3: Empirical recency versus observation period (T) for the calibration sample. Every customer necessarily lies on or below the diagonal recency = T ; the spread of points away from the diagonal encodes the latent-activity signal the BG/NBD model exploits.

3.3.3 Model Assumptions and Limitations

The BG/NBD model relies on several underlying assumptions:

- **Stationarity.** Each customer's transaction rate λ and dropout probability p are assumed constant over time. Real-world behaviour can shift owing to lifecycle progression, marketing activity, or broader market trends.
- **Independence of purchase and dropout.** The framework presumes no relationship between how frequently a customer purchases and their likelihood of churning. In practice, highly engaged customers may display distinct attrition patterns.
- **Absence of covariates.** The standard model excludes customer-level factors such as demographics, acquisition source, or marketing exposure. The hierarchical extensions of Section 3.5 partially mitigate this by allowing segment-level variation.

Despite these simplifications, the BG/NBD model has consistently performed well in empirical validation studies and remains the predominant probabilistic approach for analysing

non-contractual customer populations (Fader & Hardie, 2009).

3.4 The Gamma-Gamma Model for Monetary Value

The BG/NBD model predicts the number of future transactions but does not model their value. The Gamma-Gamma model, introduced by Fader and Hardie (2013), complements the BG/NBD by modelling the monetary value of individual transactions, enabling complete CLV estimation when combined with the transaction-count predictions.

3.4.1 Model Specification

The Gamma-Gamma model assumes the following generative process:

- **Individual spending.** Each customer's average transaction value z follows a Gamma distribution, $z \mid p, \nu \sim \text{Gamma}(p, \nu)$, where ν is individual-specific.
- **Heterogeneity across customers.** The individual-specific rate ν varies across customers according to another Gamma distribution, $\nu \sim \text{Gamma}(q, \gamma)$.

This structure produces a marginal distribution for average transaction values that accommodates the right-skewed spending distributions typically observed in retail data. The model's three parameters (p, q, γ) are estimated at the population level.

3.4.2 Key Assumption: Independence of Frequency and Monetary Value

A critical assumption of the Gamma-Gamma model is that there is no systematic relationship between a customer's purchase frequency and their average transaction value. This is necessary for the monetary and frequency components of CLV to be modelled independently and combined multiplicatively. In practice the assumption should be verified empirically by computing the correlation between purchase frequency and average transaction value in the calibration data (Figure 3.4). Violation can lead to biased CLV estimates, and the empirical analysis explicitly tests this condition.

3.4.3 CLV Computation

With both models estimated, individual-level CLV over a future horizon of t periods is computed as

$$\text{CLV}_i = \underbrace{E[\text{transactions}_i(t)]}_{\text{BG/NBD}} \times \underbrace{E[\text{monetary}_i]}_{\text{Gamma-Gamma}} \times \text{margin factor}. \quad (3.4)$$




Figure 3.4: Correlation structure among the RFM summary statistics in the calibration data. The frequency–monetary association is the quantity that must be near zero for the Gamma-Gamma independence assumption to hold.

In the Bayesian formulation, both the expected transaction count and the expected monetary value are posterior distributions rather than point estimates. CLV is therefore itself a posterior distribution, enabling probabilistic statements such as “there is a 90% probability that this customer’s 12-month CLV lies between \$150 and \$420.” This distributional CLV is the foundation for the decision-theoretic targeting framework of Section 3.6.

3.5 Hierarchical Extensions: Partial Pooling Across Segments

The standard BG/NBD and Gamma-Gamma models estimate a single set of population-level parameters for the entire customer base, implicitly assuming all customers are drawn from the same behavioural distribution. In practice, meaningful differences may exist across segments, for example, customers in different countries, acquisition channels, or product categories. Hierarchical (multilevel) Bayesian models address this by introducing segment-level parameters that are themselves drawn from a shared hyperprior. This produces *partial pooling*: segment-level estimates are pulled toward the global mean in proportion to their uncertainty, borrowing statistical strength from data-rich segments to stabilise estimates for data-sparse ones (Gelman & Hill, 2006). Figure 3.5 depicts this structure.




Figure 3.5: Hierarchical (partial-pooling) structure across country segments. Each segment’s behavioural parameters are drawn from a shared global hyperprior, so estimates for data-sparse segments (DE, FR, Other) are pulled toward the global mean, borrowing strength from the data-rich UK segment. Complete pooling would collapse the middle layer; no pooling would remove the global hyperprior entirely.

3.5.1 Partial Pooling vs. Alternatives

Three estimation strategies exist for multi-segment data, summarised in Table 3.2.

Table 3.2: Estimation strategies for multi-segment customer data.

Strategy	Approach	Limitation
Complete pooling	Ignore segment structure; estimate one global model	Ignores real differences between segments
No pooling	Estimate fully independent models per segment	Noisy estimates for small segments; no information sharing
Partial pooling	Hierarchical model; segment parameters drawn from a shared prior	More complex to implement; requires careful prior specification

The UCI Online Retail II dataset contains customers from multiple countries (UK, Germany, France, Netherlands, among others), with the UK accounting for the majority of transactions and other markets contributing substantially fewer observations (Figure 3.6). This creates an ideal testing ground for hierarchical modelling: the UK segment provides abundant data to estimate behavioural parameters precisely, while smaller country segments benefit from partial pooling that borrows strength from the global estimate. RQ2 directly tests whether this hierarchical structure improves CLV prediction for data-sparse segments.

3.5.2 Implementation Considerations

Hierarchical models introduce additional computational challenges. The *centred* parameterisation, in which segment parameters are drawn directly from the hyperprior, can create pathological posterior geometries known as Neal’s funnel (Neal, 2003). These geometries




Figure 3.6: Customer counts by country segment in the calibration data. The pronounced imbalance, a dominant UK segment alongside several small markets, is precisely the regime in which partial pooling is expected to help.

cause MCMC samplers to explore inefficiently, producing low effective sample sizes and divergent transitions. The standard solution is the *non-centred* parameterisation, which reparameterises the model so that the sampler operates on standardised (unit-normal) variables that are then scaled and shifted by the hyperparameters. This breaks the problematic correlation structure and typically resolves funnel-related sampling difficulties. The empirical implementation employs non-centred parameterisation where diagnostics indicate it is necessary.

3.6 Decision Theory Under Uncertainty

A distinctive advantage of Bayesian CLV models is that they produce full posterior distributions over customer-level CLV, enabling principled decision-making under uncertainty. This section introduces the decision-theoretic framework that RQ3 evaluates empirically.

3.6.1 The Targeting Problem

Consider a retention campaign with a fixed per-customer cost c . The firm must decide which customers to target. Under a point-estimate approach, the decision rule is typically: target all customers whose estimated CLV exceeds c . This ignores the uncertainty around the estimate; a customer with a CLV point estimate of \$100 and a cost threshold of \$90 is treated identically regardless of whether the 90% credible interval is [\$95, \$105] (high confidence) or [\$20, \$180] (low confidence).

Under a Bayesian decision-theoretic framework, the targeting decision instead uses the full posterior. The optimal decision for each customer depends on the posterior probability that their CLV exceeds the cost threshold:

$$\text{Target customer } i \quad \text{if} \quad P(\text{CLV}_i > c \mid \text{data}) > \tau, \quad (3.5)$$

where τ is a probability threshold chosen by the firm to balance the cost of false positives (targeting customers who would not generate sufficient value) against the cost of false negatives (missing valuable customers). This framework naturally accounts for estimation uncertainty and can be tuned to the firm's risk appetite. Figure 3.7 contrasts two customers with the same posterior mean but different uncertainty, illustrating why the probabilistic rule can differ from the point-estimate rule.




Figure 3.7: Two customers with the *same* posterior-mean CLV but different uncertainty. A point-estimate rule treats them identically. The decision-theoretic rule targets customer i when $P(\text{CLV}_i > c) > \tau$: customer A places almost all posterior mass above the cost c and is a confident target, whereas customer B's mass straddles c , so the probabilistic rule appropriately discounts the apparent value.

3.6.2 Expected Improvement Over Point Estimates

The decision-theoretic approach is expected to outperform point-estimate targeting in two specific scenarios:

- **High-uncertainty customers near the threshold.** Customers whose posterior CLV distribution spans the cost threshold benefit most from the probabilistic treatment. Point estimates arbitrarily classify them as above or below the threshold, while the posterior probability provides a calibrated measure of confidence.
- **Data-sparse segments.** Customers with limited purchase history have wider posterior distributions. The decision-theoretic framework naturally expresses this uncertainty, preventing overconfident targeting of customers whose apparent high value may be driven by sampling noise.

The empirical analysis tests whether these expected improvements materialise on the UCI Online Retail II holdout data.

3.7 Summary of Theoretical Findings and Hypotheses

This chapter has established the Bayesian probabilistic framework for CLV prediction. The key theoretical findings are:

- The BG/NBD + Gamma-Gamma framework provides a complete generative model for non-contractual CLV that jointly captures purchase timing, latent attrition, and monetary value, producing full posterior distributions over individual-level CLV.
- Hierarchical extensions enable partial pooling across segments, which theory predicts will improve estimates for data-sparse segments by borrowing strength from data-rich ones.
- Decision-theoretic targeting using posterior CLV distributions provides a principled framework for risk-aware customer selection that should outperform point-estimate approaches, particularly for high-uncertainty customers near the targeting threshold.

These findings convert the three gaps of Section 2.7 into testable hypotheses, collected in Table 3.3: H1 operationalises Gap 1 (calibrated Bayesian inference), H2 operationalises Gap 2 (partial pooling under segment imbalance), and H3 operationalises Gap 3 (posterior-based decisions). Chapter 4 specifies, for each hypothesis, the models, data partition, and metrics by which it is judged.

Table 3.3: Testable hypotheses evaluated in the empirical analysis.

ID	Hypothesis
H1	The Bayesian BG/NBD + Gamma-Gamma model will achieve competitive or superior out-of-sample accuracy compared to RFM-heuristic and XGBoost baselines, with the additional benefit of calibrated uncertainty estimates.
H2	Hierarchical partial pooling across country segments will reduce prediction error for small-market customers (e.g., Germany, France) compared to both complete-pooling and no-pooling alternatives.
H3	Decision-theoretic targeting using $P(\text{CLV} > c)$ will produce higher cumulative realised value than targeting based on point-estimate CLV rankings.

Chapter 4

Methodology

This chapter describes the empirical strategy used to evaluate the three hypotheses of Section 3.7. It specifies the research design, the dataset and its preparation, the Bayesian and baseline models, the inference procedure, and the evaluation protocol that links each metric to a hypothesis.

4.1 Research Design

The study adopts a *calibration–holdout* predictive-validation design, the standard protocol for customer-base analysis (Fader et al., 2005). The transaction history is split chronologically at a fixed cut-off date: data before the cut-off (the *calibration* period) are used to estimate every model, and data after it (the *holdout* period) provide the ground truth against which predictions are scored. Because the split is temporal rather than random, the design measures genuine forecasting performance and precludes information leakage from the future into model fitting.

All models are trained on identical calibration inputs and evaluated on identical holdout targets, so differences in performance are attributable to the models rather than to the data partition. The design instruments each research question with a dedicated component, completing the gap–hypothesis–method chain of Section 2.7: Gap 1/RQ1 through the accuracy and uncertainty-calibration protocol (coverage, CRPS; Section 4.9); Gap 2/RQ2 through the three-way pooling comparison and per-segment error analysis (Section 4.5); and Gap 3/RQ3 through the realised-value targeting simulation (Section 4.9).

4.2 Dataset

The empirical analysis uses the UCI *Online Retail II* dataset, a transaction-level record of a UK-based online retailer spanning 1 December 2009 to 9 December 2011. The raw data comprise 1,067,371 line items across two annual sheets, with fields for invoice number, stock code, description, quantity, unit price, invoice timestamp, customer identifier, and country. The setting is non-contractual and episodic: customers purchase sporadically with no subscription, making activity status latent and the dataset an appropriate testbed for the probabilistic models of Chapter 3.

4.3 Data Preparation

4.3.1 Cleaning

The raw transactions are cleaned in five steps, implemented in `src/data.py`: (i) rows with a missing customer identifier are dropped, since customer-level modelling is impossible without them; (ii) cancellations, identified by invoice numbers beginning with “C”, are removed; (iii) non-product stock codes representing operational charges (postage, bank charges, manual adjustments, and similar) are filtered out; (iv) rows with non-positive quantity or price are discarded; and (v) line-level revenue is computed as $\text{quantity} \times \text{price}$. Cleaning reduces the data from 1,067,371 to 802,679 line items, removing 24.8% of rows. Figure 4.1 visualises the attrition at each step.




Figure 4.1: Data-cleaning waterfall: the cumulative effect of each filtering step on the row count, from 1,067,371 raw line items to 802,679 retained transactions.

4.3.2 Temporal Split

The cleaned data are partitioned at $\text{CAL_END} = 1 \text{ March } 2011$. The calibration period (1 December 2009 to 28 February 2011, 473,386 transactions) yields a maximum customer observation window of 65.1 weeks; the holdout period (1 March 2011 to 9 December 2011, 329,293 transactions) spans 40.4 weeks and defines the prediction horizon over which forecasts are scored. Figure 4.2 shows the monthly revenue series and the location of the split.

4.3.3 Customer-Level Aggregation

Calibration transactions are collapsed to one row per customer, producing the summary statistics required by the BG/NBD and Gamma-Gamma models. Following the conventions




Figure 4.2: Monthly transaction revenue across the full observation window. The calibration/holdout boundary is fixed at 1 March 2011.

of Fader et al. (2005) and Fader and Hardie (2013):

- **Frequency** (x) is the number of *repeat* purchases, i.e. the count of purchase occasions minus one. One-time buyers have $x = 0$.
- **Recency** (t_x) is the time, in weeks, from a customer's first to their last purchase.
- T is the time, in weeks, from the first purchase to the end of the calibration period.
- **Monetary value** (m) is the mean revenue per repeat transaction; the first purchase is excluded, per the Gamma-Gamma convention, and one-time buyers are assigned $m = 0$.

Aggregation yields 4,522 customers, of whom 67.0% are repeat purchasers and 33.0% are one-time buyers. The mean repeat-purchase frequency is 3.78 (median 1, maximum 200), reflecting the heavy right skew typical of retail. Among repeat purchasers, the mean value per transaction is £384.99 (median £298.52). Time is expressed in weeks throughout (`time_unit = "W"`).

4.3.4 Country Segmentation

To support the hierarchical analysis (RQ2), customers are labelled by country. Countries with fewer than 30 customers are collapsed into an “Other” segment to avoid unstable per-segment estimates, yielding four segments: United Kingdom (4,152 customers), Other (242), Germany (74), and France (54). The pronounced imbalance (the UK accounts for over 91% of customers) is exactly the regime in which partial pooling is expected to help the small segments (Figure 3.6).

4.4 Exploratory Findings

Before modelling, the calibration sample is examined to characterise the input distributions and to test a key modelling assumption. The frequency distribution (Figure 4.3) is heavily right-skewed, with most customers making few repeat purchases and a long tail of highly active buyers. The monetary distribution (Figure 4.4) is similarly right-skewed, motivating the Gamma-Gamma model’s accommodation of skewed spend.




Figure 4.3: Distribution of repeat-purchase frequency in the calibration sample. The mass at low frequencies and the long right tail are characteristic of non-contractual retail.

Test of the Gamma-Gamma independence assumption. The Gamma-Gamma model requires that purchase frequency and average transaction value be independent (Section 3.4.2). Empirically, the correlation between frequency and monetary value among repeat purchasers is weak: Pearson $r = 0.13$ and Spearman $\rho = 0.21$ (Figure 3.4). The association is small enough that the independence assumption is approximately satisfied, but the mildly positive sign indicates that more frequent buyers spend slightly more per order; this caveat is revisited when interpreting the monetary-model results and in the limitations (Section 6.4).




Figure 4.4: Distribution of average transaction value among repeat purchasers. The right skew motivates the Gamma-Gamma specification.

4.5 Model Specification

Three Bayesian models are implemented in PyMC (`src/models.py`), following the generative structure of Chapter 3.

Pooled BG/NBD. The standard BG/NBD model (Fader et al., 2005) is estimated using the `pymc-marketing` implementation, which expresses the likelihood in the numerically stable form of Fader et al. (2005) (their expression 4). This form avoids the underflow and the ill-conditioned posterior geometry that a naive coding of the two-branch likelihood produces, and it samples efficiently. The population parameters r, α, a, b are assigned weakly informative half-normal priors whose scales are set from summary statistics of the calibration data (`src/priors.py`), which places prior mass in the region the data support and further stabilises sampling. The latent per-customer transaction rate and dropout probability are integrated out analytically, so the sampler operates only on the four population parameters.

Hierarchical BG/NBD. The hierarchical variant is a custom PyMC model that reuses the same stable likelihood but places the four BG/NBD parameters at the country-segment level, each drawn from a shared hyperprior. To avoid the funnel geometry described in Section 3.5.2, a non-centred parameterisation is used: standardised normal offsets $z \sim \text{Normal}(0, 1)$ per segment are scaled and shifted by the hyperparameters and mapped to the positive reals through an exponential (log-normal) link. A deliberately tight between-segment scale prior controls the residual funnel induced by the small country segments. This yields segment-specific (r, α, a, b) while permitting partial pooling toward the global mean.

Gamma-Gamma. The Gamma-Gamma monetary model (Fader & Hardie, 2013) is estimated with the corresponding `pymc-marketing` implementation, fitted only on repeat purchasers ($x > 0$), since average transaction value is undefined for one-time buyers. Its three population parameters receive half-normal priors, and expected spend is obtained from the model’s posterior. CLV for one-time buyers imputes the population mean spend.

4.6 Bayesian Inference

All posteriors are obtained by Hamiltonian Monte Carlo using the No-U-Turn Sampler (Section 3.1.1). Each model is sampled with 4 chains of 1000 draws following 1000 tuning iterations and a fixed random seed of 42 for reproducibility. The target acceptance probability is 0.9 for the pooled and Gamma-Gamma models and is raised to 0.95 for the hierarchical model to control its sharper posterior geometry. Convergence is assessed with the diagnostics of Section 3.1.1: $\widehat{\text{split-}\hat{R}}$ (convergence requires values below 1.01), effective sample size, the number of divergent transitions, and the Bayesian fraction of missing information. These diagnostics are reported in Section 5.1 and determine whether reparameterisation was required for the hierarchical model.

4.7 Baseline Models

Four non-Bayesian baselines (`src/baselines.py`) contextualise the probabilistic models:

- **Naive (mean).** Predicts the population-average purchase rate and average spend for every customer; a lower bound on useful performance.
- **RFM heuristic.** Scores customers into quintiles on recency, frequency, and monetary value, and projects future transactions from a composite score scaled by the base purchase rate.
- **XGBoost (two-stage).** Gradient-boosted trees trained on engineered features (core RFM, derived ratios, inter-purchase-time statistics, and day-of-week/spend-trend temporal features), with separate models for the holdout transaction count and per-transaction spend, combined multiplicatively into CLV. This mirrors the BG/NBD + Gamma-Gamma decomposition, making the Bayesian-versus-ML comparison clean.
- **Pareto/NBD (optional).** The maximum-likelihood predecessor to BG/NBD, included where the `lifetimes` library is available.

4.8 Prediction

For the Bayesian models, predictions are computed as expectations over the posterior. For each posterior draw, the conditional expected number of transactions over the holdout horizon, $E[X(t) | x, t_x, T]$, and the probability of being alive, $P(\text{alive})$, are evaluated for every customer; the Gamma-Gamma posterior yields the expected per-transaction value. Customer-level CLV is the product of the expected transaction count and expected value (Equation 3.4). Because the computation is performed per draw, every prediction is itself a posterior distribution, from which means and credible intervals are derived. For quantities concerning *realised* outcomes, namely predictive intervals and the targeting probability $P(\text{CLV} > c)$, the sampling variability of the outcome is added by drawing counts from the conditional predictive distribution per posterior draw, yielding the posterior predictive distribution; the distinction between the posterior of an expectation and the posterior predictive proves consequential in Sections 5.2.2 and 5.4.

4.9 Evaluation Protocol

Models are scored on the holdout period using a suite of metrics (`src/evaluation.py`), each mapped to a hypothesis (Table 4.1).

Table 4.1: Mapping of evaluation metrics to research hypotheses.

Hypothesis	Metrics
H1 (accuracy + uncertainty)	MAE, RMSE, MAPE for transactions and CLV; Spearman and Gini for ranking; NDCG@ k ; credible-interval coverage and mean interval width; CRPS; $P(\text{alive})$ AUC and Brier score; decile calibration tables
H2 (hierarchical pooling)	Per-country MAE for each model; posterior shrinkage of segment parameters toward the global mean
H3 (decision-theoretic targeting)	Cumulative realised net value under point-estimate, posterior-probability, and oracle targeting, swept over targeting depth and intervention cost; targeting lift

Because several of these metrics are less familiar than MAE, Table 4.2 gives a one-line reading of what each captures and why it is relevant here, before the formal definitions that follow.

Accuracy and ranking. Mean absolute error and root mean squared error measure point accuracy for both holdout transactions and CLV; the Gini coefficient and Spearman corre-

Table 4.2: Evaluation metrics at a glance: plain-language intuition and why each is used.

Metric	What it measures / why it matters here
MAE, RMSE	Average size of the prediction error (RMSE penalises large misses more); lower is better
MAPE	Error as a percentage of the true value; readable but unstable when truths are near zero
Spearman	Whether the model <i>rank</i> s customers in the right order, ignoring the exact values
Gini	Concentration of value the ranking captures; how well high-value customers rise to the top
NDCG@ k	Quality of the <i>top k</i> of the ranking specifically; the metric most aligned with targeting a limited budget
Coverage	Share of realised outcomes that fall inside the stated interval; should match the nominal level (e.g. 90%)
Interval width	Sharpness of the intervals; narrower is better <i>only</i> if coverage is maintained
CRPS	Single score rewarding a predictive distribution that is both accurate and well calibrated; lower is better
AUC, Brier	Quality of $P(\text{alive})$ as an activity classifier: ranking (AUC) and probability calibration (Brier)

lation measure ranking quality, and NDCG@100 and NDCG@500 measure top- k ranking relevant to targeting.

Uncertainty calibration. For the Bayesian models, credible-interval coverage compares the nominal and empirical coverage of posterior predictive intervals (well-calibrated intervals achieve coverage close to nominal), the mean interval width quantifies sharpness, and the continuous ranked probability score (CRPS) jointly rewards accuracy and calibration. These metrics operationalise the “calibrated uncertainty” claim of H1, which point-prediction baselines cannot satisfy.

Activity classification. $P(\text{alive})$ is evaluated as a binary classifier of holdout activity (a customer is active if they make at least one holdout purchase; 59.0% of customers, 2,670 of 4,522, are active) using AUC and the Brier score.

Decision-theoretic targeting. For H3, a targeting simulation ranks customers by (i) posterior mean CLV, (ii) the probability $P(\text{CLV} > c)$ computed from the posterior *predictive* distribution of realised CLV, and (iii) realised value (oracle), and reports the cumulative net value captured at targeting depths of 5–50%. The intervention cost c is swept over

$\{\pounds100, \pounds300, \pounds600, \pounds1,200, \pounds2,000\}$, a grid scaled to the magnitude of realised customer value in the holdout, to ensure the conclusion is robust to the cost assumption; $\pounds600$ is the headline case.

4.10 Implementation and Reproducibility

The pipeline is implemented in Python using PyMC 5 for inference (Salvatier et al., 2016), the `pymc-marketing` library for the standard BG/NBD and Gamma-Gamma implementations (PyMC Labs, 2024), ArviZ for posterior handling and diagnostics, scikit-learn and XGBoost for baselines, and Matplotlib for figures. The full analysis is orchestrated by `src/run_all_models.py`, which executes data preparation, MCMC sampling, prediction, evaluation, table generation, and plotting as a single reproducible sequence. All stochastic components use fixed seeds, posterior traces are serialised to NetCDF, and result tables are exported as both CSV and $\text{\LaTeX}$ for direct inclusion in this document.

Chapter 5

Results

This chapter reports the empirical results of the calibration–holdout evaluation and tests the three hypotheses of Section 3.7. Section 5.1 establishes that the MCMC inference is trustworthy; Sections 5.2–5.4 address H1, H2, and H3 in turn; and Section 5.5 collates the verdicts.

Note on reproduction. The tables and figures in this chapter are produced by `src/run_all_models.py` and written to `outputs/results/` and `outputs/figures/`; all numeric values quoted in the prose are read directly from those generated tables. The standard BG/NBD and Gamma-Gamma models are estimated with the `pymc-marketing` implementation of the numerically stable Fader–Hardie likelihood, while the hierarchical model is a custom PyMC implementation of the same likelihood with partial pooling across country segments.

5.1 MCMC Convergence and Diagnostics

Before interpreting any posterior, the quality of the sampling is verified using the diagnostics of Section 3.1.1. Figure 5.1 shows the split- $\hat{R}$ values for the pooled BG/NBD, hierarchical BG/NBD, and Gamma-Gamma models. Convergence is indicated when all parameters satisfy $\hat{R} < 1.01$; the effective sample sizes and the number of divergent transitions for each model are summarised below. The trace and pair plots (Figures 5.2 and 5.3) provide a visual check on mixing and posterior correlation structure.




Figure 5.1: Split- $\hat{R}$ convergence diagnostics for the pooled BG/NBD model. The dashed line marks the 1.01 threshold.




Figure 5.2: Posterior trace plots for the pooled BG/NBD population parameters (r, α, a, b) . Well-mixed, stationary chains indicate reliable sampling.




Figure 5.3: Pairwise posterior marginals for the BG/NBD parameters, with divergences highlighted. Strong correlations would motivate reparameterisation.

For the hierarchical model, a non-centred parameterisation (Section 3.5.2) with a tight between-segment scale was adopted to control the funnel geometry that the small country segments induce; Figure 5.4 confirms that this resolved the sampling pathologies. Across all three models the maximum split- $\hat{R}$ is 1.002 and the minimum bulk effective sample size exceeds 1,100, with no divergent transitions. The diagnostics therefore indicate that the reported posteriors are well converged and provide a sound basis for the inferences that follow.

5.2 H1: Predictive Accuracy and Calibrated Uncertainty

5.2.1 Point and Ranking Accuracy

Table 5.1 reports the headline comparison across all models for both holdout transaction prediction and CLV, covering error (MAE, RMSE, MAPE), ranking (Spearman, Gini), and top- k relevance (NDCG). H1 predicts that the Bayesian BG/NBD + Gamma-Gamma model is competitive with or superior to the RFM-heuristic and XGBoost baselines. The results show that the Bayesian model attains a transaction MAE of 1.74 against 2.12 for XGBoost and 2.33 for the RFM heuristic, and a CLV Gini of 0.86 (versus 0.80 for XGBoost). The ranking metrics are the most decisive: the Bayesian model reaches an NDCG@100 of 0.91 against 0.53 for XGBoost, so it identifies the highest-value customers far more reliably than either baseline while also carrying the lowest CLV MAE (£847 against £1,478 for XGBoost). The size of this ranking gap warrants explanation, since a value near 0.9 against one near 0.5 is a large margin. Two factors combine to produce it. First, the generative model borrows strength across the customer base, so even customers with sparse histories receive a sensibly ordered value estimate, whereas the discriminative XGBoost, trained to minimise average error, is comparatively indifferent to the exact ordering within the small high-value tail that NDCG@100 rewards. Second, the leakage-safe design gives XGBoost a shorter effective training window (Section 6.4.2), which handicaps precisely the top-of-ranking task; its accuracy on bulk error metrics (transaction MAE 2.12) is much closer to the Bayesian model than the NDCG gap alone would suggest. The Bayesian model is thus superior, not merely competitive, on out-of-sample accuracy.

Table 5.1: Model comparison across transaction and CLV prediction metrics.

Model	Tx MAE	Tx RMSE	Tx MAPE (%)	Tx Spearman	Tx Gini	CLV MAE (£)	CLV RMSE (£)	CLV Gini	NDCG@100	NDCG@500
Hierarchical BG/NBD	1.742	3.310	91.8	0.609	0.776	847.5	3,307.2	0.855	0.905	0.903
BG/NBD (Bayesian)	1.743	3.306	92.0	0.609	0.776	847.1	3,300.0	0.855	0.905	0.903
XGBoost (two-stage)	2.117	3.768	127.4	0.577	0.752	1,478.4	5,900.8	0.801	0.534	0.615
RFM Heuristic	2.330	6.366	58.6	0.582	0.702	1,160.5	6,728.6	0.785	0.531	0.617
Naive (mean)	3.114	6.212	186.8	–	0.052	1,619.6	6,854.3	0.070	0.017	0.061

The decile calibration plot (Figure 5.5) shows predicted versus actual transactions for each model. Points close to the diagonal indicate good calibration of the point predictions; sys-




Figure 5.4: Split- $\hat{R}$ for the hierarchical BG/NBD model under the non-centred parameterisation.

tematic departure in the top deciles would signal over- or under-prediction for the most active customers.




Figure 5.5: Decile calibration: mean predicted versus mean actual holdout transactions for each model. The diagonal denotes perfect calibration.

5.2.2 Uncertainty Calibration

The distinctive claim of H1 is that the Bayesian model adds *calibrated uncertainty* that point-prediction baselines cannot provide. Table 5.2 reports the empirical coverage of posterior predictive credible intervals, the mean interval width, and the CRPS. Coverage close to the nominal level (e.g. empirical coverage near 90% for the 90% interval) indicates well-calibrated uncertainty; the observed coverage is 89.3%, almost exactly the nominal 90%, so the posterior predictive intervals are well calibrated. The posterior predictive distributions for a sample of customers (Figure 5.6) illustrate how the model expresses heterogeneous uncertainty across the customer base.

Table 5.2: Posterior predictive credible interval coverage and CRPS.

	Coverage (90%)	Mean width	CRPS
Model			
BG/NBD (Bayesian)	0.893	4.586	1.223
Hierarchical BG/NBD	0.894	4.588	1.222




Figure 5.6: Posterior predictive distributions of holdout transactions for a sample of customers (violins), with realised values overlaid. The spread of each violin is the model’s per-customer uncertainty.

5.2.3 Activity Prediction

The $P(\text{alive})$ estimates are evaluated as a classifier of holdout activity (Table 5.3). Because a customer with no repeat purchase is trivially assigned $P(\text{alive}) = 1$ under the BG/NBD model (they cannot have dropped out after a repeat that never occurred), the discrimination metrics are computed on repeat purchasers, for whom the estimate is informative. On this population the BG/NBD model achieves an AUC of 0.71 and a Brier score of 0.21, indicating moderate but useful discrimination of customers who remain active. An AUC of 0.71 should be read as a genuine signal rather than a disappointing one: activity is a latent state that no model can observe, and a customer who is truly still active may simply not purchase again within the finite holdout window through ordinary randomness, which caps how well *any* classifier can separate the two groups on this task. The value is well above the chance level of 0.5, and, as the calibration-first framing of this thesis stresses, the probability is more useful for being calibrated than for being sharp. Figure 5.7 shows that $P(\text{alive})$ increases with calibration frequency, as theory predicts.

Table 5.3: $P(\text{alive})$ evaluation: binary activity prediction metrics.

Model	Accuracy	Precision	Recall	F1	AUC-ROC	AUC-PR	Brier
BG/NBD (Bayesian)	0.729	0.734	0.975	0.837	0.711	0.868	0.207
Hierarchical BG/NBD	0.729	0.734	0.974	0.837	0.709	0.867	0.207




Figure 5.7: Distribution of $P(\text{alive})$ overall and by frequency group. Higher-frequency customers concentrate at higher $P(\text{alive})$.

5.2.4 A Worked Example: One Customer End to End

The aggregate metrics are best understood by watching the model reason about a single customer. Customer 14817, introduced in Section 3.2.4, enters with only four numbers: frequency $x = 4$, recency $t_x = 33.1$ weeks, observation window $T = 49.2$ weeks, and mean repeat spend $m = \text{£}526$. From these the fitted model produces, in sequence, every quantity this thesis is concerned with.

First, activity. Despite the fifteen-week silence before the calibration cut-off, the pattern of four repeat purchases spread across the window yields $P(\text{alive}) = 0.88$: probably, but not certainly, still active. Second, future purchasing. The conditional expectation over the 40-week holdout horizon is $E[X] = 2.81$ transactions (90% credible interval $[2.74, 2.88]$). Third, spend. The Gamma-Gamma model shrinks the customer's observed £526 toward the population mean, giving an expected value of roughly £485 per transaction. Multiplying the two components gives an expected CLV of £1364, with a strikingly tight point-posterior interval of $[\text{£}1326, \text{£}1401]$.

That tight interval is where the worked example becomes instructive, because it reflects only *parameter* uncertainty and is therefore the wrong interval for a decision about a realised outcome. The posterior *predictive* distribution of the customer's actual holdout value, which additionally accounts for the sampling variability of the transaction count and the possibility of dropout, is far wider: a 90% interval of $[\text{£}0, \text{£}2888]$, with substantial mass at zero (Figure 5.8). The two distributions disagree sharply about the very question a targeting rule asks. At the headline intervention cost of £600, the probability $P(\text{CLV} > \text{£}600)$ computed from the posterior of the expectation is a falsely emphatic 1.00, because the narrow spike sits entirely above the threshold; computed from the predictive distribution it is a candid 0.77, acknowledging the roughly one-in-four chance the customer returns too little to justify the spend. This single customer previews the population-wide artefact examined in Section 5.4. In the event, customer 14817 returned twice in the holdout and spent £1010: profitable to target, above the cost, and comfortably inside the predictive interval that the point estimate's tight band would have hidden.

Verdict on H1. Taken together, the accuracy metrics (Table 5.1) and the uncertainty diagnostics (Table 5.2) indicate that H1 is supported: the Bayesian model is superior on accuracy, most clearly on the ranking metrics that matter for targeting, while uniquely providing well-calibrated predictive intervals.

5.3 H2: Hierarchical Partial Pooling Across Segments

H2 predicts that hierarchical partial pooling reduces prediction error for small country segments relative to the pooled (complete-pooling) model. The natural benchmark for this claim is the three-way comparison of *complete* pooling, *partial* pooling, and *no* pooling, the last being an independent BG/NBD fitted separately on each segment. Table 5.4 reports per-country transaction MAE for all three specifications alongside the classical baselines, and Figure 5.9 visualises the comparison.

Two contrasts matter. Against complete pooling, the hierarchical model does not reduce error




Figure 5.8: Two posteriors for one customer (ID 14817). The narrow posterior of $E[CLV]$ (orange) reflects parameter uncertainty only and lies entirely above the £600 cost, giving $P(CLV > £600) = 1.00$. The posterior predictive of realised CLV (blue) is far wider, placing $\sim 23\%$ of its mass below the cost, so the same probability computed correctly is 0.77. The realised holdout value (£1010) falls well inside the predictive distribution.

in the small markets: for Germany ($n = 74$) its MAE is 1.635 versus 1.622, and for France ($n = 54$) 2.318 versus 2.223, in both cases a marginal increase. Against no pooling, however, the pattern is exactly the one partial pooling is designed to produce. No pooling is the *worst* of the three in every small segment, where the sparse per-segment data over-fit (France 2.444, Germany 1.670), whereas complete pooling is instead worst in the larger, behaviourally distinct “Other” segment (1.695), where imposing the global parameters ignores a genuine difference. Partial pooling is never the worst in any segment and attains the lowest error aggregated over the three non-UK segments (1.746, against 1.755 for no pooling and 1.757 for complete pooling). For the data-rich UK segment all three coincide (≈ 1.74), as expected where data are abundant. The differences among the specifications are small in absolute terms, a consequence of the tight between-segment prior needed for stable sampling, which shrinks the segment parameters strongly toward the global values.

Table 5.4: Per-country transaction-prediction MAE by model (RQ2/H2).

Country	n	MAE: BG/NBD (Bayesian)	MAE: Hierarchical BG/NBD	MAE: Naïve (mean)	MAE: RFM Heuristic	MAE: XGBoost (two-stage)	MAE: No pooling (per segment)
United Kingdom	4,152	1.741	1.741	3.079	2.291	2.114	1.740
Other	242	1.695	1.653	3.618	2.499	2.057	1.627
Germany	74	1.622	1.635	3.136	3.031	2.366	1.670
France	54	2.223	2.318	3.529	3.586	2.329	2.444

The mechanism behind any improvement is visible in the shrinkage of the segment-level parameters toward the global mean. Figure 5.10 shows the per-segment posteriors for the BG/NBD shape parameter r : small segments exhibit wider posteriors pulled toward the global estimate, whereas the UK segment is estimated precisely and close to its no-pooling




Figure 5.9: Per-country transaction MAE by model. Differences concentrated in the small segments (Germany, France, Other) test the partial-pooling hypothesis.

value.




Figure 5.10: Per-segment posterior of the BG/NBD parameter r (90% HDI), with the global mean marked. Small segments are shrunk toward the global estimate.

Verdict on H2. H2 is partially supported. Partial pooling does not beat complete pooling on this relatively homogeneous country segmentation, but the borrowing-of-strength mechanism is demonstrably operating: it strictly dominates no pooling in every small segment, is never the worst specification anywhere, and attains the lowest error aggregated across the non-UK segments. It thus delivers the robustness that motivates the hierarchical model, adapting between the two pooling extremes. Two caveats qualify this reading: the between-

specification differences are small in absolute terms, and the tight between-segment prior required for stable sampling itself limits how far partial pooling can depart from complete pooling (Section 6.4).

5.4 H3: Decision-Theoretic Targeting

H3 predicts that targeting customers by the posterior probability $P(\text{CLV} > c)$ yields higher cumulative realised value than targeting by point-estimate CLV. Table 5.5 reports the targeting simulation for the Bayesian model: cumulative net value under the point-estimate rule, the posterior-probability rule, and the oracle, across targeting depths and the swept intervention cost. Figure 5.11 visualises the strategy curves and the incremental advantage of the posterior rule.

Table 5.5: Decision-theoretic targeting (BG/NBD (Bayesian)): posterior-probability vs point-estimate vs oracle net value, swept over intervention cost.

Cost (£)	Depth	n tgt.	Point (£)	Posterior (£)	Oracle (£)	Random (£)	Improvement (£)	Improvement (%)
100	0.05	227	2,921,236.3	2,542,570.7	3,319,869.2	268,880.2	-378,665.5	-13.0
	0.10	453	3,596,944.3	3,186,256.8	4,042,655.8	536,576.0	-410,687.5	-11.4
	0.15	679	4,033,697.4	3,698,961.8	4,505,583.1	804,271.7	-334,735.6	-8.3
	0.20	905	4,338,665.8	4,009,033.4	4,831,417.1	1,071,967.5	-329,632.4	-7.6
	0.30	1,357	4,686,448.3	4,437,284.0	5,223,076.8	1,607,358.9	-249,164.3	-5.3
	0.40	1,809	4,911,945.0	4,754,007.4	5,428,485.4	2,142,750.4	-157,937.6	-3.2
	0.50	2,261	5,053,868.3	4,966,972.2	5,524,777.4	2,678,141.9	-86,896.1	-1.7
300	0.05	227	2,875,836.3	2,596,186.3	3,274,469.2	223,480.2	-279,650.0	-9.7
	0.10	453	3,506,344.3	3,235,056.3	3,952,055.8	445,976.0	-271,288.0	-7.7
	0.15	679	3,897,897.4	3,657,604.1	4,369,783.1	668,471.7	-240,293.3	-6.2
	0.20	905	4,157,665.8	3,970,587.1	4,650,417.1	890,967.5	-187,078.7	-4.5
	0.30	1,357	4,415,048.3	4,301,850.8	4,951,676.8	1,335,958.9	-113,197.6	-2.6
	0.40	1,809	4,550,145.0	4,485,572.3	5,066,685.4	1,780,950.4	-64,572.7	-1.4
	0.50	2,261	4,601,668.3	4,581,183.2	5,072,577.4	2,225,941.9	-20,485.1	-0.4
600	0.05	227	2,807,736.3	2,603,112.1	3,206,369.2	155,380.2	-204,624.2	-7.3
	0.10	453	3,370,444.3	3,242,151.6	3,816,155.8	310,076.0	-128,292.8	-3.8
	0.15	679	3,694,197.4	3,549,658.8	4,166,083.1	464,771.7	-144,538.6	-3.9
	0.20	905	3,886,165.8	3,820,895.8	4,378,917.1	619,467.5	-65,270.0	-1.7
	0.30	1,357	4,007,948.4	3,989,952.9	4,544,576.8	928,858.9	-17,995.5	-0.4
	0.40	1,809	4,007,445.0	3,978,165.9	4,523,985.4	1,238,250.4	-29,279.1	-0.7
	0.50	2,261	3,923,368.3	3,927,762.3	4,394,277.4	1,547,641.9	4,394.0	0.1
1,200	0.05	227	2,671,536.3	2,644,226.4	3,070,169.2	19,180.2	-27,309.9	-1.0
	0.10	453	3,098,644.3	3,063,604.7	3,544,355.8	38,276.0	-35,039.6	-1.1
	0.15	679	3,286,797.4	3,253,129.5	3,758,683.1	57,371.7	-33,667.9	-1.0
	0.20	905	3,343,165.8	3,327,674.8	3,835,917.1	76,467.5	-15,491.0	-0.5
	0.30	1,357	3,193,748.4	3,195,590.4	3,730,376.8	114,658.9	1,842.0	0.1
	0.40	1,809	2,922,045.0	2,928,928.9	3,438,585.4	152,850.4	6,883.9	0.2
	0.50	2,261	2,566,768.3	2,574,752.8	3,037,677.4	191,041.9	7,984.5	0.3
2,000	0.05	227	2,489,936.2	2,487,013.3	2,888,569.2	-162,419.8	-2,922.9	-0.1
	0.10	453	2,736,244.3	2,728,504.0	3,181,955.8	-324,124.0	-7,740.3	-0.3
	0.15	679	2,743,597.4	2,742,195.7	3,215,483.1	-485,828.3	-1,401.7	-0.1
	0.20	905	2,619,165.8	2,610,630.8	3,111,917.1	-647,532.5	-8,535.0	-0.3
	0.30	1,357	2,108,148.4	2,123,555.5	2,644,776.8	-970,941.1	15,407.1	0.7
	0.40	1,809	1,474,845.0	1,481,137.2	1,991,385.4	-1,294,349.6	6,292.2	0.4
	0.50	2,261	757,968.3	767,207.2	1,228,877.4	-1,617,758.1	9,238.8	1.2

The probability score is computed from the posterior *predictive* distribution of realised CLV




Figure 5.11: Decision-theoretic targeting for the Bayesian model at the headline cost of £600. Left: cumulative net value by strategy and depth. Right: the net-value difference between the posterior-probability rule and the point estimate.

rather than from the posterior of *expected* CLV. This distinction, the same one that governs interval coverage in Section 5.2.2, is decisive here: the posterior of the expectation reflects only parameter uncertainty and, on a dataset of this size, places $P(\text{CLV} > c)$ at essentially 0 or 1 for almost every customer, so the ranking degenerates into arbitrary tie-breaking and the rule spuriously appears to fail by more than 50%. Using the predictive distribution removes that artefact. Figure 5.12 shows the effect across the whole customer base: computed from the posterior of the expectation, 96% of customers are pinned at a probability of 0 or 1, leaving almost nothing for the ranking to sort on; computed from the posterior predictive, the same probability is graded across the full $[0, 1]$ range, and the worked example of Section 5.2.4 (Figure 5.8) is one customer drawn from this second, informative distribution.




Figure 5.12: Distribution of the per-customer targeting probability $P(\text{CLV} > £600)$ across all 4,522 customers. (a) Computed from the posterior of $E[\text{CLV}]$, the probability saturates: 96% of customers sit at 0 or 1. (b) Computed from the posterior predictive of realised CLV, it is graded and usable as a decision score. This is the saturation artefact that makes the expectation-based rule fail for mechanical reasons.

Table 5.6 makes the contrast concrete on the risk-sensitive metrics that the probability rule is meant to serve: at the headline cost and depth, the predictive probability rule matches the

point estimate on both the hit rate (the share of targeted customers whose realised value exceeds the cost) and wasted spend (the cumulative shortfall on unprofitable targets), whereas the expected-CLV version of the same rule collapses on both.

Table 5.6: Risk-sensitive targeting metrics at the headline cost of £600 and a targeting depth of 20%. *Hit rate* is the fraction of targeted customers whose realised holdout value exceeds the intervention cost; *wasted spend* is the cumulative shortfall on targeted customers who fall below it. The point-estimate rule and the posterior-*predictive* probability rule are close on both; the probability rule computed from the posterior of *expected* CLV instead collapses, illustrating the saturation artefact.

Targeting rule	Hit rate	Wasted spend (£)
Point estimate, $E[CLV]$	0.848	55,040
$P(CLV > c)$, expected-CLV posterior	0.619	144,738
$P(CLV > c)$, posterior predictive	0.827	60,452

At the headline intervention cost of £600, the posterior-probability rule then captures £3.82M in cumulative net value at a targeting depth of 20%, against £3.89M for the point-estimate rule, a shortfall of only 1.7%. Across the swept cost grid (£100 to £2,000) and across targeting depths the two rules are, in practice, equivalent: the gap is within a few percentage points everywhere and turns marginally in the probability rule's favour at the highest costs and depths, where mis-targeting a customer whose realised value falls below the intervention cost becomes costly enough for its risk-aversion to matter. This near-parity is the expected outcome on reflection: ranking by posterior mean CLV is close to optimal for maximising *total* realised value, so the additional tail information the probability rule exploits is not needed for that objective. For completeness, Figure 5.13 reports the conventional targeting-lift curves, which measure revenue captured as a function of targeting depth.




Figure 5.13: Cumulative-gain and lift curves for all models. Steeper early rise indicates better identification of high-value customers.

Verdict on H3. The simulation indicates that H3 is not supported: targeting by $P(CLV > c)$ does not outperform point-estimate targeting for total realised value; the two are effectively

equivalent. The finding is informative rather than merely negative, on two counts. First, it shows that the decision value of the full posterior depends on the objective: for maximising expected captured value the posterior mean is sufficient, whereas the probability rule expresses a risk-averse preference that would be advantageous under a loss-sensitive objective or a hard budget constraint (Section 3.6.2), neither of which the total-value criterion rewards. Second, it exposes a concrete pitfall in the decision-theoretic use of Bayesian models: a rule defined on the posterior of an expectation rather than on the posterior predictive can fail for purely mechanical reasons, and respecting that distinction is essential to using the posterior correctly.

5.5 Summary of Hypothesis Tests

Table 5.7 collates the verdicts. The detailed interpretation and its implications for theory and practice are discussed in Chapter 6.

Table 5.7: Summary of hypothesis-test outcomes.

ID	Hypothesis (abridged)	Verdict
H1	Bayesian model competitive/superior on accuracy with calibrated uncertainty	Supported
H2	Partial pooling reduces error for small country segments	Partially supported
H3	$P(\text{CLV} > c)$ targeting beats point-estimate targeting	Not supported

Chapter 6

Discussion and Conclusion

This final chapter synthesises the empirical findings, situates them within the theoretical framework of Chapters 2 and 3, and draws out the contributions, limitations, and avenues for further work.

6.1 Summary of Findings

The thesis set out to evaluate a Bayesian probabilistic approach to CLV prediction in a non-contractual e-commerce setting, organised around three research questions. The calibration–holdout evaluation on the UCI Online Retail II dataset (4,522 customers, a 40.4-week holdout horizon) yields the following picture.

RQ1 / H1: Accuracy with calibrated uncertainty. The Bayesian BG/NBD + Gamma-Gamma model was found to be superior to the RFM-heuristic and XGBoost baselines on out-of-sample accuracy, most clearly on the ranking metrics that matter for targeting (Table 5.1), while uniquely providing posterior predictive intervals whose empirical coverage was well calibrated at 89.3% against a nominal 90% (Table 5.2). The central argument of the thesis, that the value of the Bayesian approach lies not only in point accuracy but in the trustworthy quantification of uncertainty, is therefore borne out by the evidence.

RQ2 / H2: Hierarchical partial pooling. H2 was partially supported (Table 5.4). Partial pooling did not improve on complete pooling for the data-sparse markets, but the borrowing-of-strength mechanism was clearly at work: no pooling was the worst specification in every small segment, complete pooling was worst in the one behaviourally distinct segment, and partial pooling was never the worst and had the lowest error aggregated across the non-UK segments. The hierarchical model thus delivered robustness rather than outright improvement, on a country segmentation that proved too homogeneous for segment-specific parameters to pay off.

RQ3 / H3: Decision-theoretic targeting. H3 was not supported: across the swept intervention costs, targeting by the posterior predictive probability $P(\text{CLV} > c)$ and by point-estimate CLV were in practice equivalent (Table 5.5). Two lessons follow. First, the decision value of the posterior depends on the objective: for maximising total captured value the

posterior mean is close to optimal, while the probability rule encodes a risk-averse preference that the total-value criterion does not reward. Second, the analysis surfaced a practical pitfall: the probability score must be computed from the posterior *predictive*, not from the posterior of expected CLV, which saturates at 0 or 1 and makes the rule appear to fail for purely mechanical reasons.

6.2 Theoretical Contributions

The thesis makes three connected contributions, each answering one of the gaps of Section 2.7. First, it provides an end-to-end Bayesian implementation of the BG/NBD + Gamma-Gamma framework estimated by Hamiltonian Monte Carlo, rather than the maximum-likelihood fitting that predominates in practice, and evaluates it calibration-first: the demonstration that posterior predictive coverage matches its nominal level on real transaction data is itself part of the contribution, since calibration is asserted far more often than it is tested. Second, it tests the hierarchical extension against *both* pooling extremes, showing that the value of partial pooling on imbalanced customer bases is robustness (never the worst specification) rather than uniform error reduction, a more precise characterisation than the literature’s usual two-way comparison supports. Third, it operationalises Bayesian decision theory for targeting and, in diagnosing why the naive rule fails, isolates a general methodological pitfall: decision rules computed from the posterior of an *expectation* saturate and degrade, and must instead use the posterior predictive distribution. This distinction is rarely made explicit in the applied CLV literature and transfers directly to the recent covariate-rich and neural extensions of customer-base analysis (Bachmann et al., 2021; Valendin et al., 2022), which produce distributional outputs facing exactly the same hazard.

6.3 Managerial Implications

For practitioners, the results carry three practical messages. First, where decisions are risk-sensitive (deciding how much to spend retaining a customer whose value is uncertain), a model that reports calibrated intervals is more useful than one that reports only a number, even when their point accuracy is similar; here the Bayesian model was additionally the more accurate. Second, firms operating across heterogeneous markets or channels can stabilise estimates for their smaller segments through partial pooling, which the results show is the safe default: never the worst option, and best in aggregate across small segments. Third, the targeting results temper expectations about probability-based rules: for maximising total campaign value, ranking by expected CLV suffices, and firms should reserve $P(\text{CLV} > c)$ rules for genuinely loss-sensitive objectives or hard budget constraints. Whichever rule is used, it must be computed from the posterior predictive distribution; computing it from the

posterior of expected CLV collapsed the hit rate from 0.85 to 0.62 and nearly tripled wasted spend in the simulation (Table 5.6).

6.4 Limitations

The limitations fall into three groups: assumptions of the models, choices in the evaluation design, and the scope of the CLV construct as measured.

6.4.1 Modelling Assumptions

The BG/NBD model assumes stationary transaction and dropout rates and independence between purchasing and dropout (Section 3.3.3); real customers change behaviour over time and in response to marketing, and the decile calibration analysis indeed shows systematic over-prediction of 20–30% for moderately active customers, consistent with mid-frequency cohorts churning faster than the stationary model allows. The Gamma-Gamma independence assumption is mildly violated (frequency–monetary Pearson $r = 0.13$, Spearman $\rho = 0.21$; Section 4.4), which may bias the monetary component slightly. Spend for the 1,493 one-time buyers (33.0% of the customer base) is imputed at the population mean, so within this substantial group the CLV ranking is driven entirely by the transaction model and spend heterogeneity is invisible. Finally, the posterior predictive counts are generated by a conditional-mean Poisson approximation that conditions away the alive/dead mixture; the resulting intervals are nevertheless empirically well calibrated (89.3% at nominal 90%), which justifies the approximation on this dataset without guaranteeing it elsewhere.

6.4.2 Evaluation Design

The evaluation uses a single temporal calibration/holdout split, so all reported differences are estimates from one realisation of the data, without formal statistical tests; the large H1 margins (transaction MAE 1.74 against 2.12) are unlikely to be artefacts of the split, but the small between-specification differences in the H2 comparison, which reach only the third decimal for the UK segment, should be read as descriptive patterns rather than statistically established effects. A paired resampling procedure over customers would put uncertainty bounds on these model differences and is a natural extension. Relatedly, the tight between-segment prior scale adopted for sampling stability itself restricts how far partial pooling can depart from complete pooling, so the H2 conclusion is conditional on that prior choice; a prior-sensitivity analysis would establish its robustness. The XGBoost baseline required an inner temporal split to avoid target leakage, which gives it a shorter effective training window than the Bayesian models; its comparatively weak top-of-ranking performance should be

interpreted with that asymmetry in mind.

6.4.3 Scope of the CLV Construct

The quantity validated here is expected revenue over the 40.4-week holdout horizon, computed with a unit margin factor and no discounting; it is a finite-horizon revenue forecast used as a CLV proxy, since only a finite window can ever be validated against realised outcomes, rather than a lifetime discounted-cash-flow estimate. Cancellations were removed during cleaning rather than netted against prior purchases, slightly overstating realised revenue for customers who returned goods, and the models omit customer-level covariates (acquisition channel, demographics, marketing exposure) that could sharpen predictions.

6.5 Generalisability

The external validity of the findings is bounded by the empirical setting in three ways. First, the evidence derives from a single UK-based giftware retailer observed over 2009–2011; transfer to other product categories, price points, purchase cadences, and market conditions remains to be established. Second, the customer population mixes consumer and wholesale buyers, with the most active customer making 200 repeat purchases in 65 weeks; such procurement-driven purchasing violates the BG/NBD’s behavioural story, and the model’s strong performance despite this heterogeneity is reassuring but not guaranteed to replicate in purer B2B settings. Segmenting by customer type rather than by geography is a promising alternative that might also provide a sharper test of H2, since the country segments examined here proved behaviourally homogeneous. Third, the holdout window includes the pre-Christmas trading peak, so the forecast horizon is seasonally favourable for a retailer of this kind; a holdout spanning a demand trough could yield different calibration. Against these bounds stand two considerations that support broader applicability: the modelling inputs are the minimal recency/frequency/monetary summaries available to any transaction-recording firm, and the methodological findings, calibration-first evaluation and the expectation-versus-predictive distinction in decision rules, are properties of the Bayesian workflow rather than of this dataset.

6.6 Future Work

These limitations suggest a clear research agenda. Relaxing stationarity through time-varying or seasonal extensions would better capture lifecycle and calendar effects. Incorporating covariates, via a regression layer on the BG/NBD and Gamma-Gamma parameters, would let the model exploit the rich behavioural data available in e-commerce

while retaining its generative, uncertainty-aware structure. Replacing the assumed independence of frequency and value with a jointly specified model would address the mild violation observed here. Methodologically, formal Bayesian model comparison (WAIC, LOO-CV; Section 3.1) across the pooled, hierarchical, and covariate-augmented specifications would put model selection on a principled footing. Finally, validating the decision-theoretic targeting rule in a live A/B test, rather than a holdout simulation, would close the loop between posterior inference and realised business value.

6.7 Concluding Remarks

Customer lifetime value in non-contractual settings is fundamentally a problem of reasoning about an unobservable state under uncertainty. This thesis has argued, and set out to demonstrate empirically, that a Bayesian treatment of generative customer-base models is well matched to that problem: it delivers competitive predictions, calibrated uncertainty, principled pooling across heterogeneous segments, and decisions that account for what the firm does not know. The broader lesson is that, for value-based marketing decisions, how confident a model is can matter as much as what it predicts, and a framework that quantifies both is the more dependable foundation for action.

References

- Angelopoulos, A. N., & Bates, S. (2023). Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4), 494–591.
- Bachmann, P., Meierer, M., & Näf, J. (2021). The role of time-varying contextual factors in latent attrition models for customer base analysis. *Marketing Science*, 40(4), 783–809.
- Berger, J. O. (1985). *Statistical decision theory and bayesian analysis* (2nd ed.). Springer.
- Blattberg, R. C., & Deighton, J. (1996). Manage marketing by the customer equity test. *Harvard Business Review*, 74(4), 136–144.
- Blattberg, R. C., Kim, B.-D., & Neslin, S. A. (2008). *Database marketing: Analyzing and managing customers*. Springer.
- Chamberlain, B. P., Cardoso, A., Liu, C. H. B., Pagliari, R., & Deisenroth, M. P. (2017). Customer lifetime value prediction using embeddings. *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1753–1762.
- Cullinan, G. J. (1977). *Picking them by their batting averages' recency-frequency-monetary method of controlling circulation* (tech. rep. No. Manual Release 2103). Direct Mail/Marketing Association.
- Dwyer, F. R. (1989). Customer lifetime valuation to support marketing decision making. *Journal of Direct Marketing*, 3(4), 8–15.
- Dwyer, F. R. (1997). Customer lifetime valuation to support marketing decision making. *Journal of Direct Marketing*, 11(4), 6–13.
- Efron, B., & Morris, C. (1975). Data analysis using stein's estimator and its generalizations. *Journal of the American Statistical Association*, 70(350), 311–319.
- Fader, P. S., & Hardie, B. G. S. (2009). Probability models for customer-base analysis. *Journal of Interactive Marketing*, 23(1), 61–69.
- Fader, P. S., & Hardie, B. G. S. (2013). *The gamma-gamma model of monetary value* (tech. rep.) (Technical note). University of Pennsylvania.
- Fader, P. S., Hardie, B. G. S., & Lee, K. L. (2005). “Counting Your Customers” the Easy Way: An Alternative to the Pareto/NBD Model. *Marketing Science*, 24(2), 275–284.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., & Rubin, D. B. (2013). *Bayesian data analysis* (3rd ed.). CRC Press.
- Gelman, A., & Hill, J. (2006). *Data analysis using regression and multilevel/hierarchical models*. Cambridge University Press.

- Gupta, S., Hanssens, D., Hardie, B., Kahn, W., Kumar, V., Lin, N., Ravishanker, N., & Sri-ram, S. (2006). Modeling customer lifetime value. *Journal of Service Research*, 9(2), 139–155.
- Gupta, S., & Lehmann, D. R. (2003). Customers as assets. *Journal of Interactive Marketing*, 17(1), 9–24.
- Hoffman, M. D., & Gelman, A. (2014). The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research*, 15, 1593–1623.
- Hughes, A. M. (1994). *Strategic database marketing*. Probus Publishing.
- Hughes, A. M. (2011). *Strategic database marketing* (4th ed.). McGraw-Hill.
- Lovallo, D., & Kahneman, D. (2003). Delusions of success: How optimism undermines executives' decisions. *Harvard Business Review*, 81(7), 56–63.
- McCarthy, D. M., & Fader, P. S. (2018). Customer-based corporate valuation for publicly traded noncontractual firms. *Journal of Marketing Research*, 55(5), 617–635.
- Neal, R. M. (2003). Slice sampling. *The Annals of Statistics*, 31(3), 705–767.
- PyMC Labs. (2024). *Pymc-marketing: Bayesian marketing analytics in python* (Version 0.x). <https://github.com/pymc-labs/pymc-marketing>
- Reichheld, F. F., & Sasser, W. E. (1990). Zero defections: Quality comes to services. *Harvard Business Review*, 68(5), 105–111.
- Reinartz, W. J., & Kumar, V. (2000). On the profitability of long-life customers in a noncontractual setting: An empirical investigation and implications for marketing. *Journal of Marketing*, 64(4), 17–35.
- Rossi, P. E., Allenby, G. M., & McCulloch, R. (2005). *Bayesian statistics and marketing*. John Wiley & Sons.
- Rust, R. T., Zeithaml, V. A., & Lemon, K. N. (2000). *Driving customer equity: How customer lifetime value is reshaping corporate strategy*. Free Press.
- Salvatier, J., Wiecki, T. V., & Fonnesbeck, C. (2016). Probabilistic programming in python using pymc3. *PeerJ Computer Science*, 2, e55.
- Schmittlein, D. C., Morrison, D. G., & Colombo, R. (1987). Counting your customers: Who are they and what will they do next? *Management Science*, 33(1), 1–24.
- Stone, B. (1995). *Successful direct marketing methods* (5th ed.). NTC Business Books.
- Storbacka, K. (1997). Segmentation based on customer profitability: Retrospective analysis of retail bank customer bases. *Journal of Marketing Management*, 13(5), 479–492.
- Valendin, J., Reutterer, T., Platzer, M., & Kalcher, K. (2022). Customer base analysis with recurrent neural networks. *International Journal of Research in Marketing*, 39(4), 988–1018.
- Vehtari, A., Gelman, A., & Gabry, J. (2017). Practical bayesian model evaluation using leave-one-out cross-validation and waic. *Statistics and Computing*, 27(5), 1413–1432.

- Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., & Bürkner, P.-C. (2021). Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing convergence of mcmc. *Bayesian Analysis*, 16(2), 667–718.
- Venkatesan, R., & Kumar, V. (2004). A customer lifetime value framework for customer selection and resource allocation strategy. *Journal of Marketing*, 68(4), 106–125.